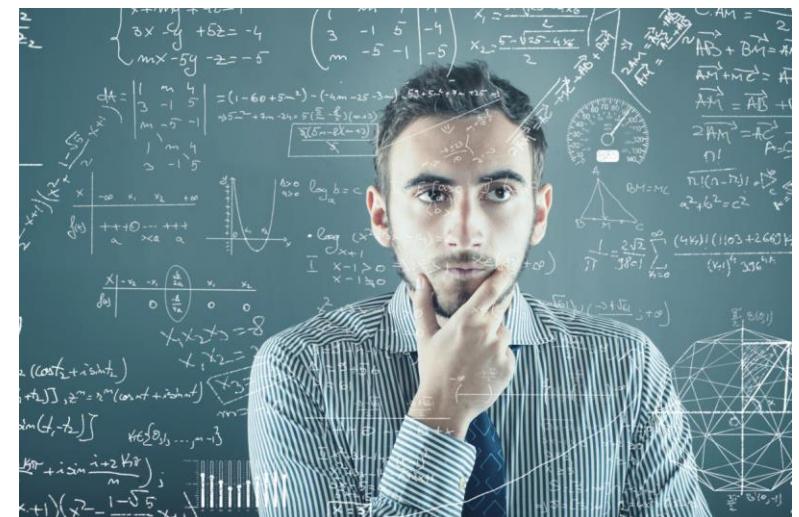
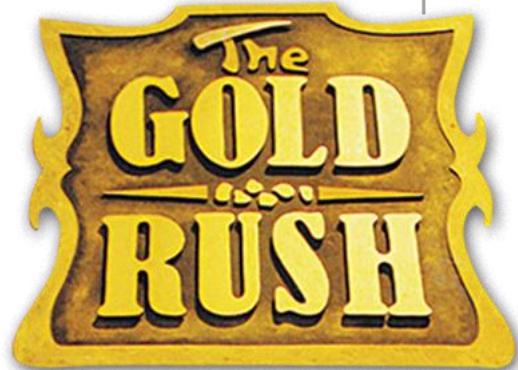


# Data Science and Exploration using Big- data Tools

Ahmed Eldawy  
Computer Science and Engineering

# Data Science

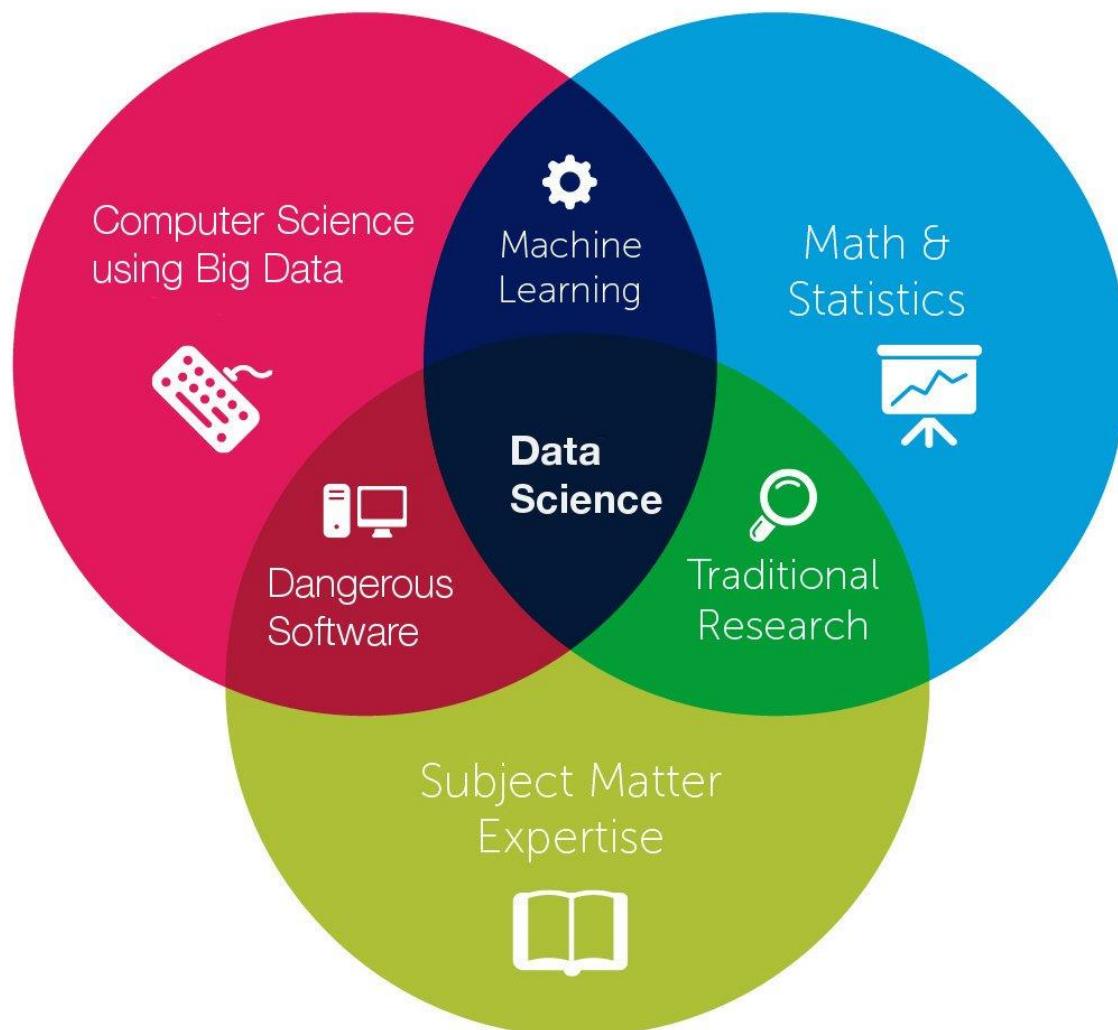
- › Data is the new gold
- › Data science is the power of extracting meaningful knowledge from data
- › Becoming increasingly important due to the abundance of data



# Data Collection



# Data Science



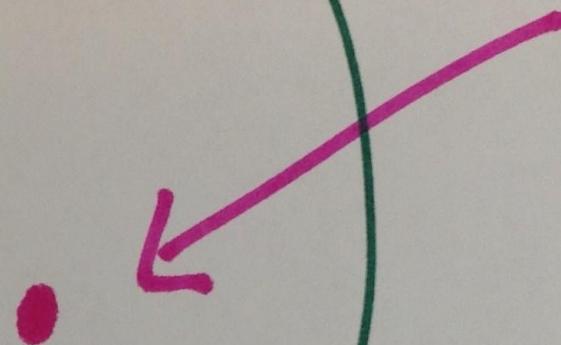
**Big Data**

Straight Ahead



All of the  
information

Information  
you  
need!



# The Market of Big Data



≡     **Forbes** / Tech / #BigData

Your roadmap for distributed care is here. Read the eBook. CISCO

JAN 20, 2017 @ 09:27 AM 72,834 ▶ The Little Black Book of Billionaire Secrets

## 6 Predictions For The \$203 Billion Big Data Analytics Market

**Gil Press, CONTRIBUTOR**  
I write about technology, entrepreneurs and innovation. [FULL BIO](#) ▾  
Opinions expressed by Forbes Contributors are their own.

*Shutterstock*

The creation and consumption of data continues to grow by leaps and bounds and with it the investment in big data

RELATED KEYWORDS ▶

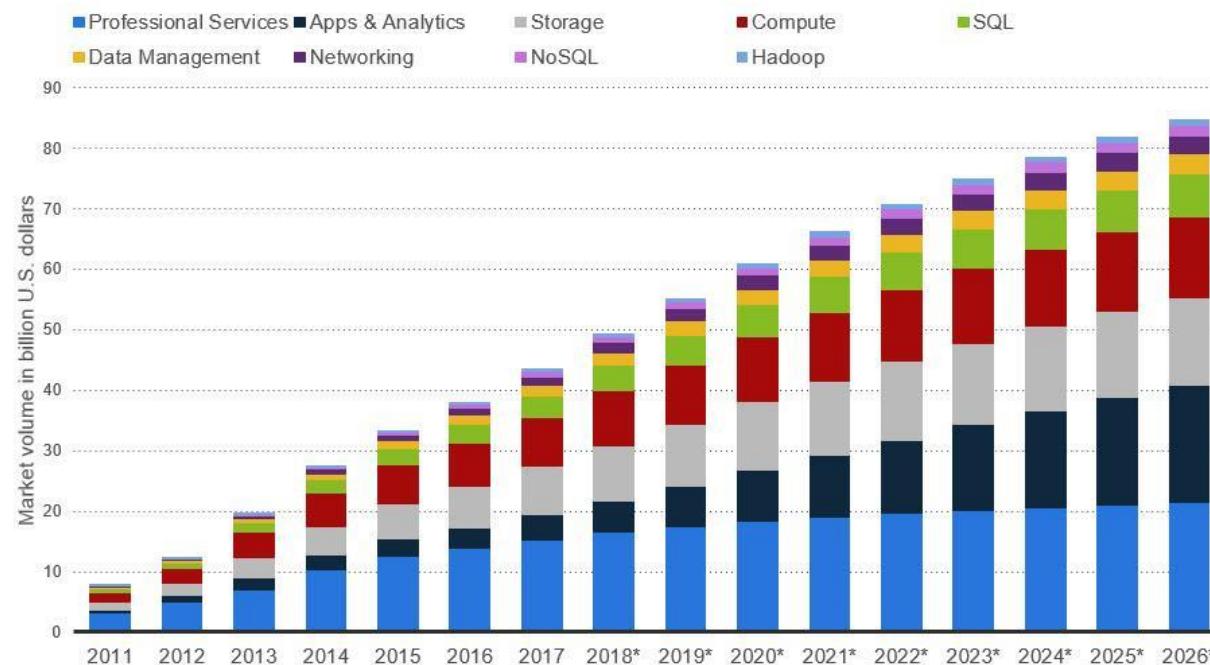
1. [BIG DATA ANALYTIC TOOLS](#) >
2. [BIG DATA ANALYTIC TRENDS](#) >
3. [BIG DATA TRENDS FOR 2018](#) >
4. [BIG DATA FOR BUSINESS](#) >
5. [NEW BIG DATA SOLUTIONS](#) >
6. [DATA MANAGEMENT PLATFORM](#) >
7. [DATA ENTRY SERVICES](#) >
8. [DATA ANALYTICS TRAINING](#) >
9. [BIG DATA COURSES](#) >
10. [GEOSPATIAL DATA MANAGEMENT](#) >

# Job Market

UCR

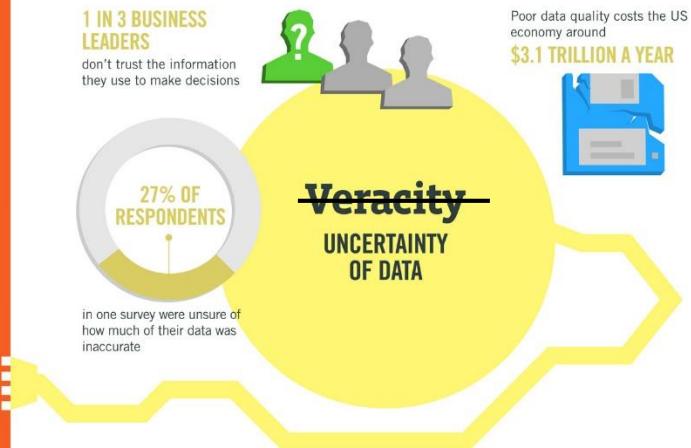
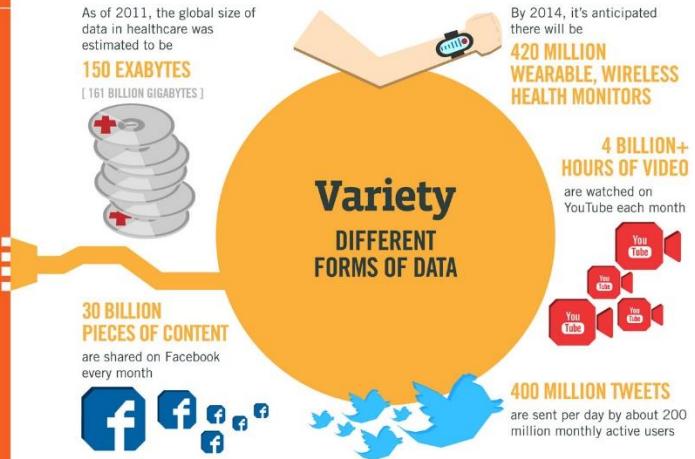
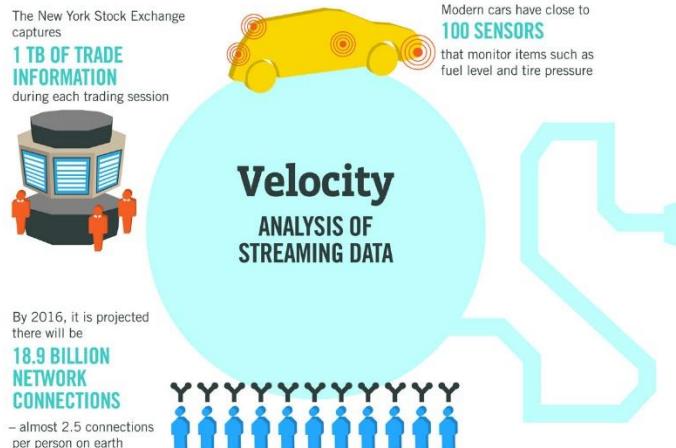
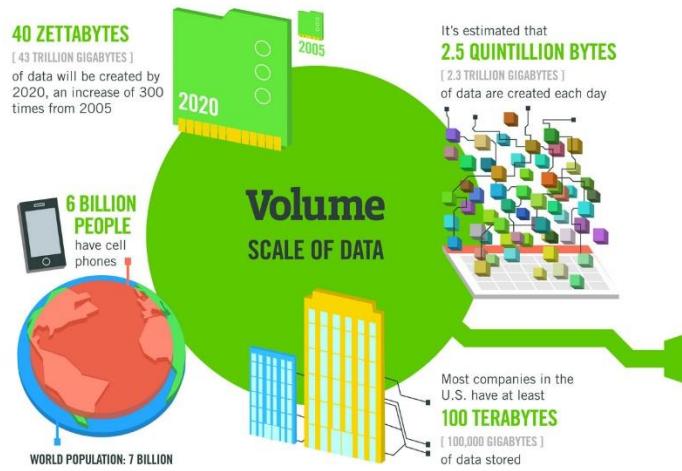
Big Data Market Worldwide Segment Revenue Forecast 2011-2026

**Big Data Market Forecast Worldwide from 2011 to 2026, by segment  
(in billion U.S. dollars)**



statista

# Four Three V's of Big Data



Sources: McKinsey Global Institute, Twitter, Cisco, Gartner, EMC, SAS, IBM, MEPTEC, QAS

# Big Data Vs Big Computation



- › Full scans (e.g., log processing)
- › Range scans
- › Point lookups
- › Iterations
- › Joins (self, binary, or multiway)
- › Proximity queries
- › Closures and graph traversals
- › **Big Data:** The computation is too expensive to be processed using the common tools

# Landscape of Big Data



- What is currently known as big data?
  - Distributed computing
  - Horizontal scaling
  - Fault tolerance
  - In-situ computation.
- No one-size fits all
  - There are many solutions and your choice depends on what you need
- You need to have basic knowledge that allows you to explore new big data systems

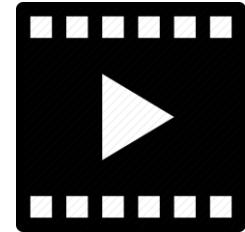
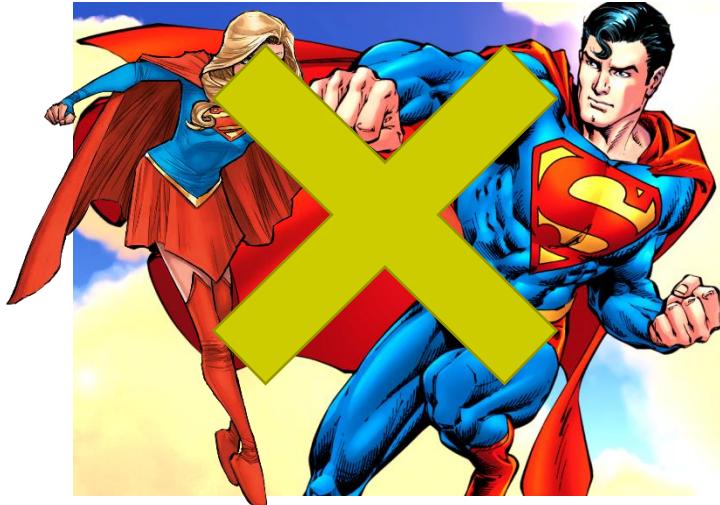
# Big-data Expert



- › Understand how the big-data platforms really work
- › Control those thousands of processors efficiently to carry out your task



# Superhero



# Ant-Man/Wasp

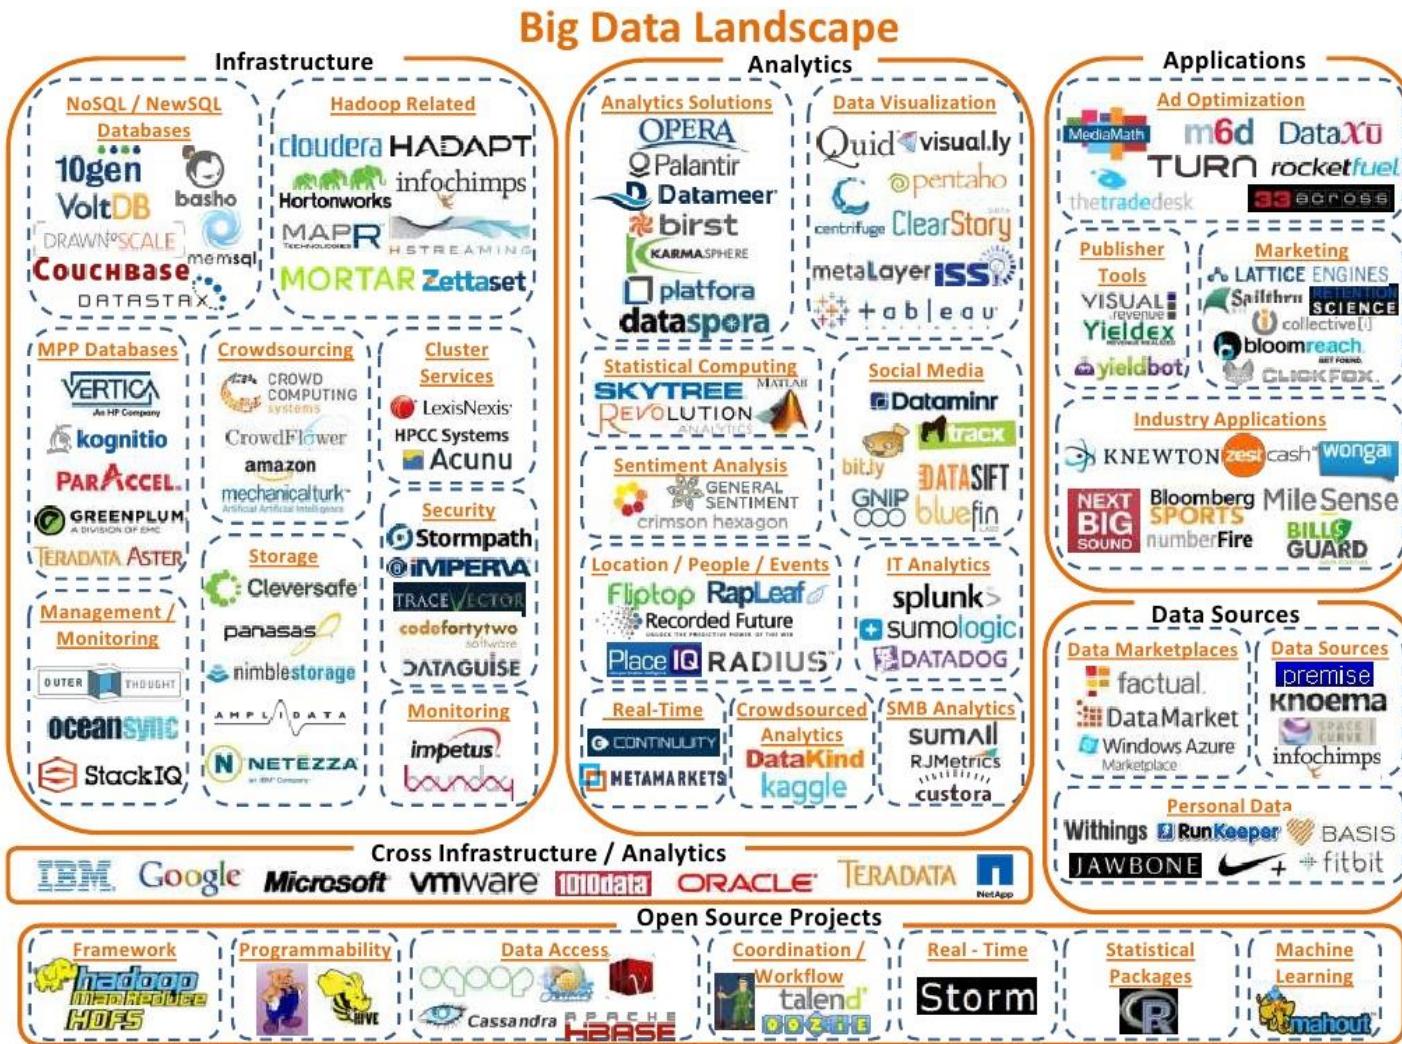


Get smaller to understand how ants work and what they are capable of.

Use this knowledge to control thousands of ants and do amazing things!

# Big Data Landscape 2012

UCR



© Matt Turck (@mattturck) and ShivonZilis (@shivonz)

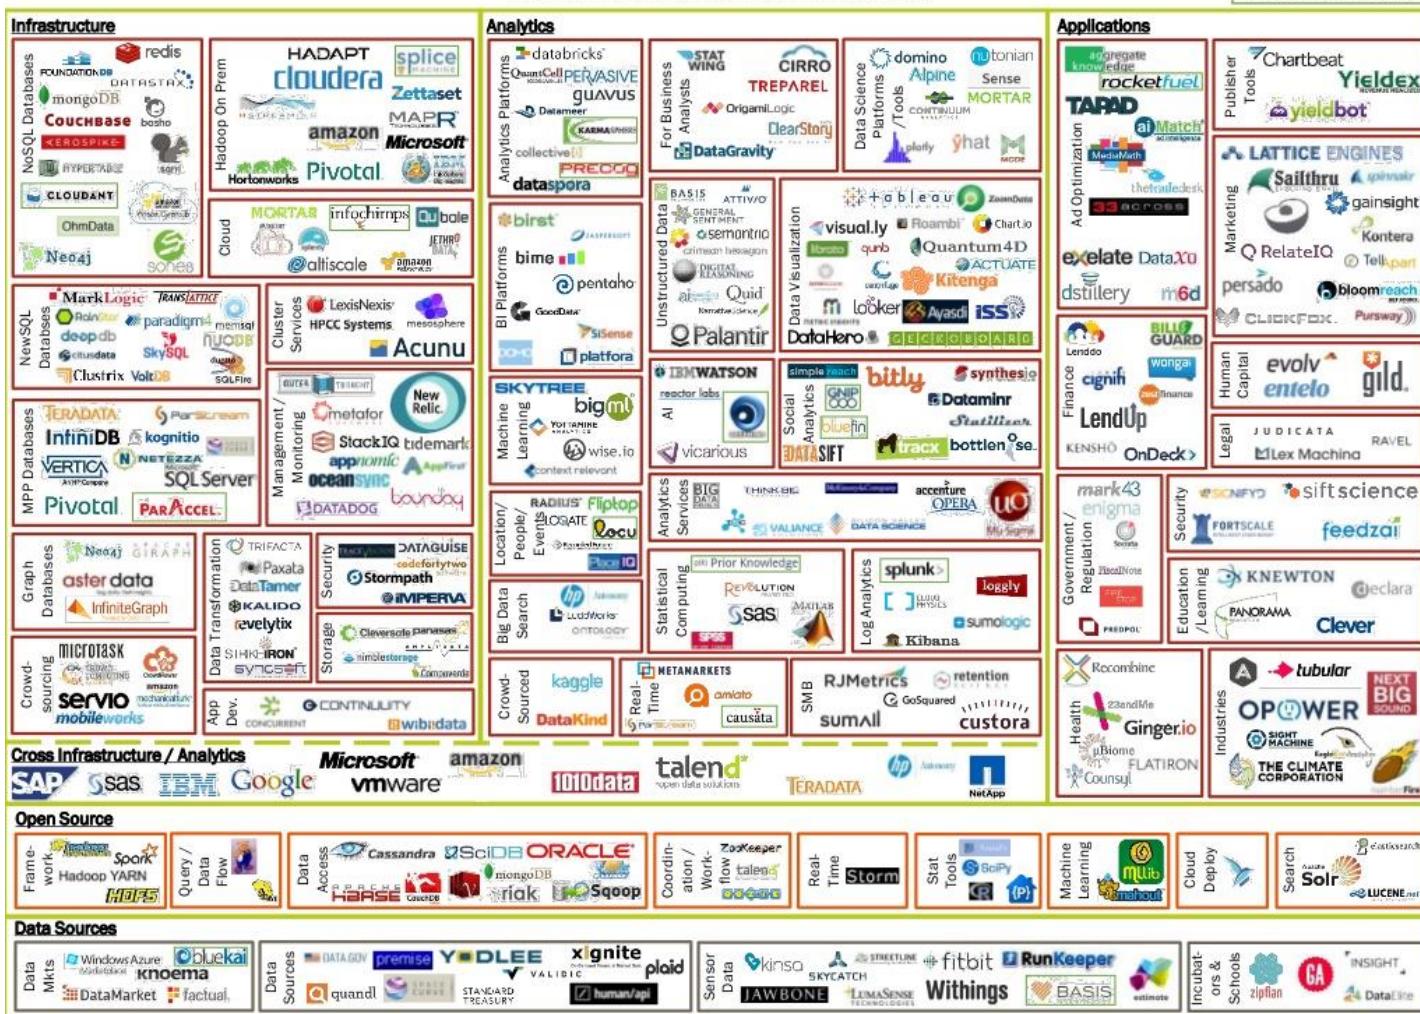
<http://mattturck.com/2012/06/29/a-chart-of-the-big-data-ecosystem/>

# Big Data Landscape 2014



## **BIG DATA LANDSCAPE, VERSION 3.0**

Exited: Acquisition or IPO



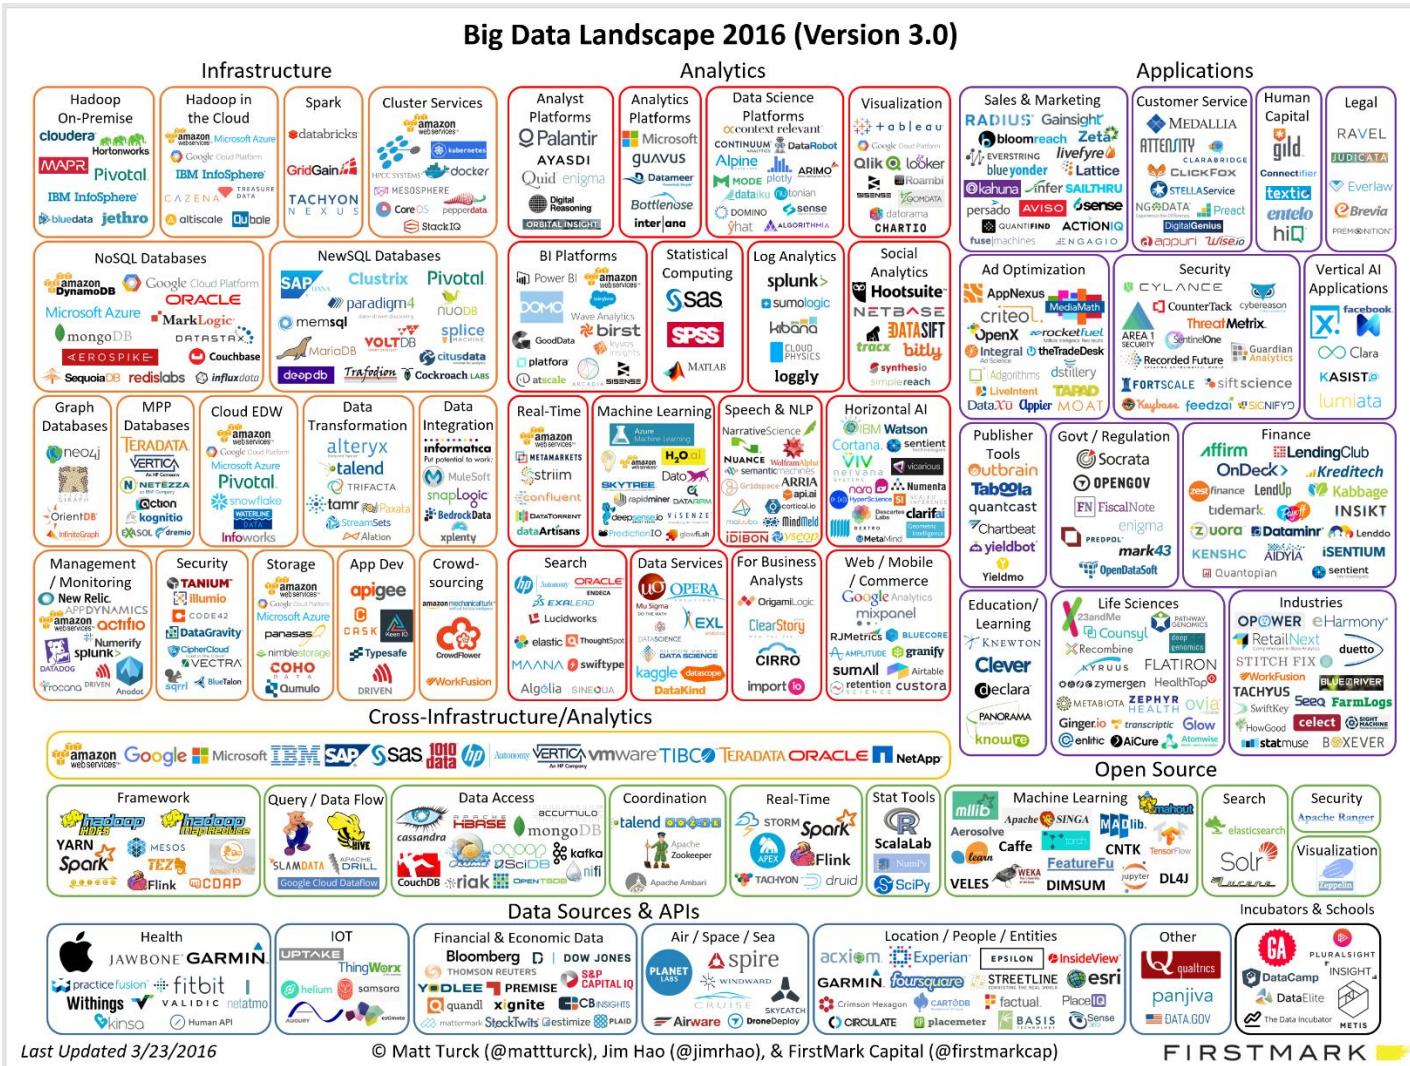
© Matt Turck (@mattturck), Sutian Dong (@sutiandong) & FirstMark Capital (@firstmarkcap)

<http://mattturck.com/2014/05/11/the-state-of-big-data-in-2014-a-chart/>

# Big Data Landscape 2016



**Big Data Landscape 2016 (Version 3.0)**



Last Updated 3/23/2016

© Matt Turck (@mattturck), Jim Hao (@jimrhao), & FirstMark Capital (@firstmarkcap)

FIRSTMARK

<http://mattturck.com/2016/02/01/big-data-landscape/>

# Big Data Landscape 2018



BIG DATA & AI LANDSCAPE 2018



V1 – Last updated 6/19/2018

© Matt Turck (@mattturck), Demilade Obavomi (@demi\_ obavomi), & FirstMark (@firstmarkcap) matturck.com/bigdata2018

FIRSTMARK  
EARLY STAGE VENTURE CAPITAL

# Components of Big Data



## Big-data Libraries

MLlib (Machine Learning), GraphX

## High-level Languages

SparkSQL, Pig, SQL++, HiveQL

## Distributed Computing

MapReduce (Hadoop and Google), Resilient Distributed Dataset (Spark), Hyracks (AsterixDB)

## Big Data Distributed Storage

Hadoop Distributed File System, Cloud storage systems (Amazon S3 and Google File System), Key-value stores

## Cloud Services

Amazon Web Services, Microsoft Azure, and Google Cloud Platform

## Coordination/Cluster Management

Oozie, Yarn, Kubernetes

# Big-data Storage



Big-data Libraries

High-level Languages

Distributed Computing

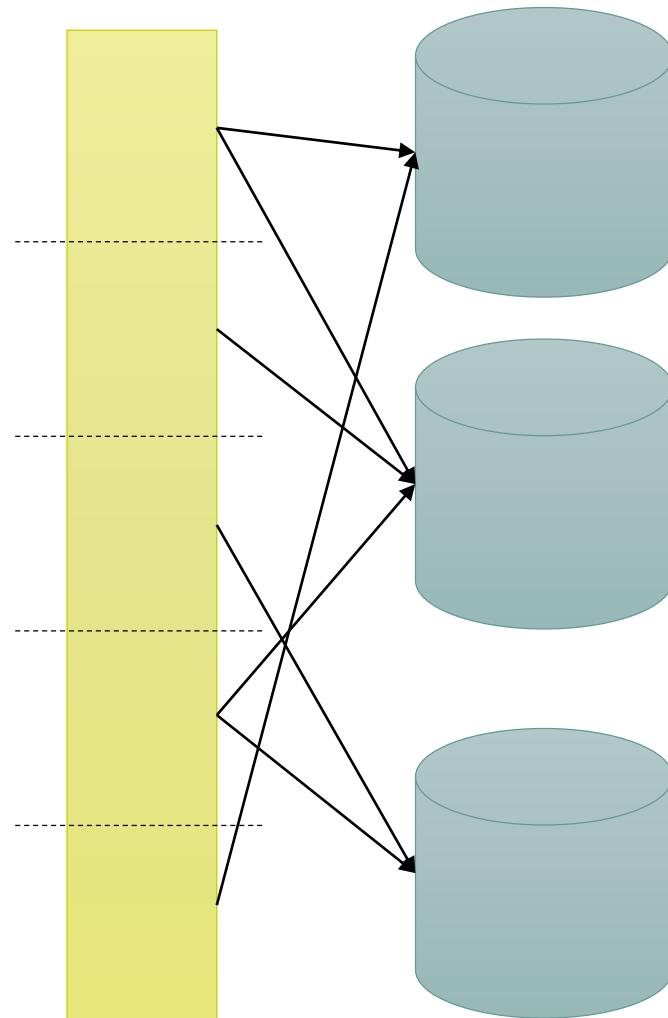
Big Data Storage

Cloud Services

Coordination/  
Cluster  
Management

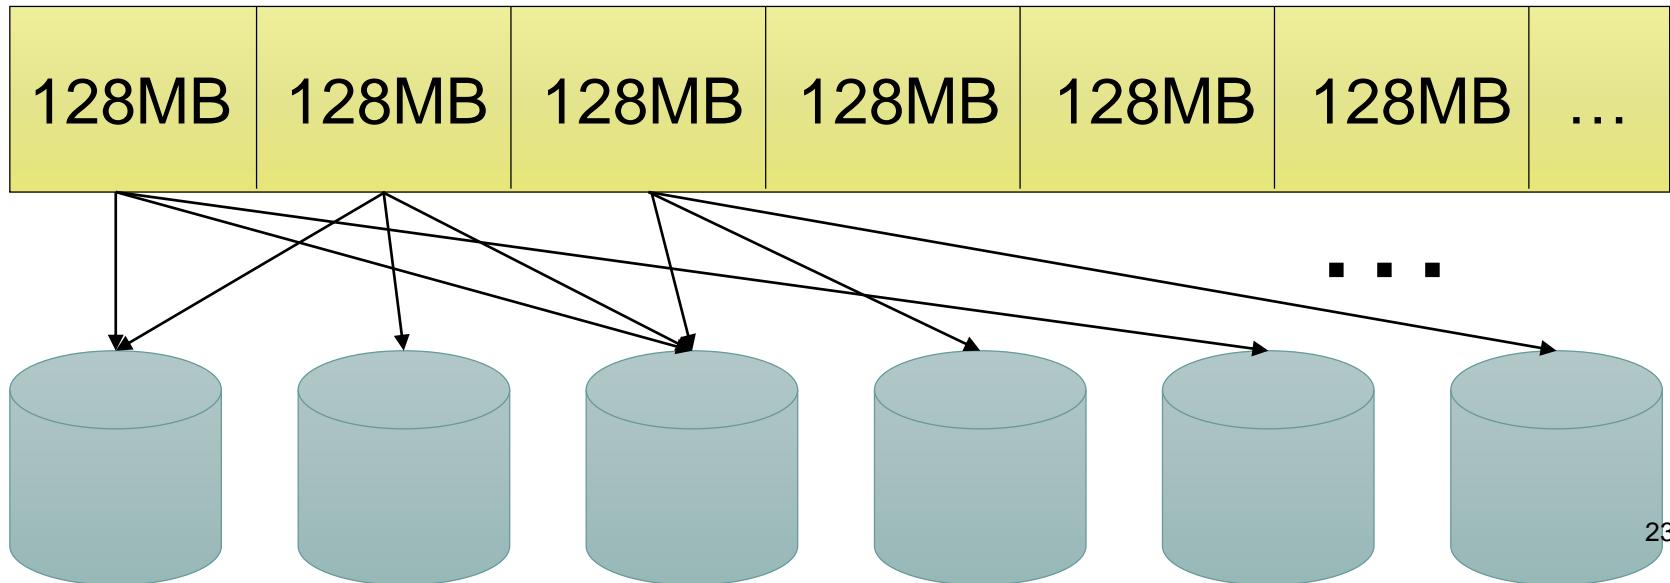
# Storage of Big Data

- › Data is growing faster than Moore's Law
- › Too much data to fit on a single machine
- › Partitioning
- › Replication
- › Fault-tolerance



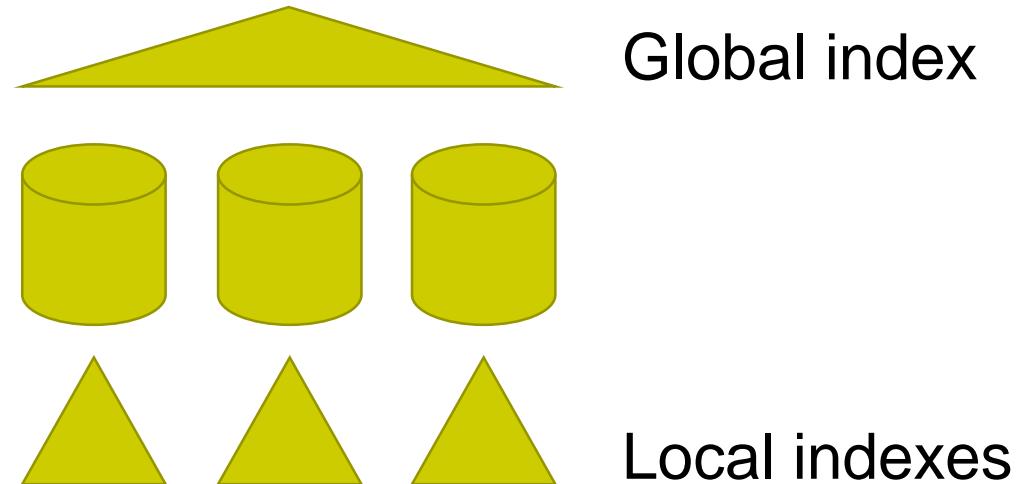
# Hadoop Distributed File System (HDFS)

- › The most widely used distributed file system
- › Fixed-sized partitioning
- › 3-way replication
- › Write-once read-many
- › See also: GFS, Amazon S3, Azure Blob Store



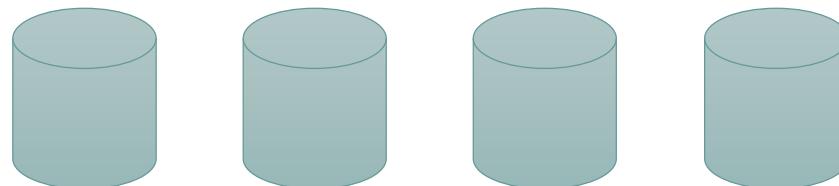
# Indexing

- › Data-aware organization
- › Global Index **partitions** the records into blocks
- › Local Indexes organize the records in a partition
- › Challenges:
  - › Big volume
  - › HDFS limitation
  - › New programming paradigms
  - › Ad-hoc indexes

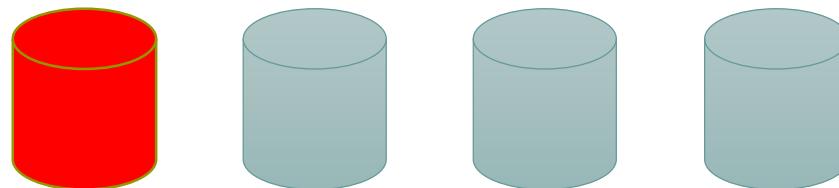


# Fault Tolerance

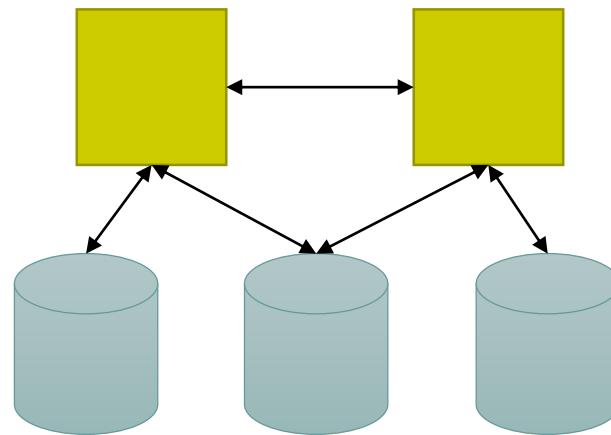
- Replication



- Redundancy

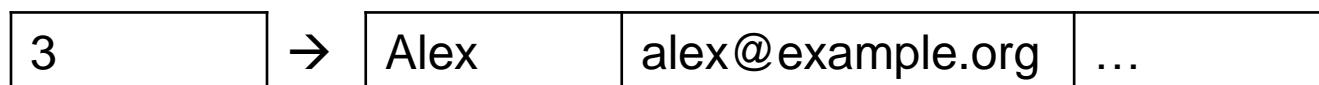
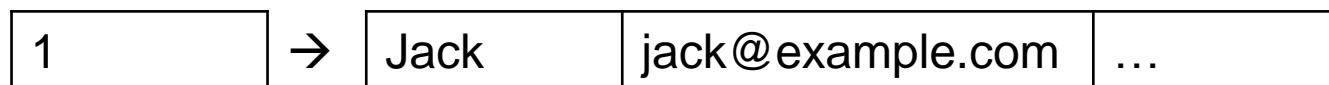


- Multiple masters



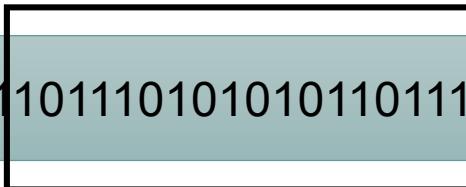
# Key-value Stores

ID	Name	Email	...
1	Jack	jack@example.com	
2	Jill	jill@example.net	
3	Alex	alex@example.org	



# Streaming

```
...10001000101010111011101010110111010111011101110100...
```



- Sub-second latency for queries
- One scan over the data
- (Partial) preprocessing
- Continuous queries
- Eviction strategies
- In-memory indexes

# Structured/Semi-structured



ID	Name	Email	...
1	Jack	jack@example.com	
2	Jill	jill@example.net	
3	Alex	alex@example.org	

## Document 1

```
{ "id": 1, "name":"Jack", "email":  
"jack@example.com", "address": {"st  
"900 university ave", "city": "Riverside",  
"CA"}, "friend_ids": [3, 55, 123]}
```

## Document 2

```
{ "id": 2, "name": "Jill", "email":  
"jill@example.net", "hobbies": ["hiking",  
"cooking"]}
```

# Distributed Computing



Big-data Libraries

High-level Languages

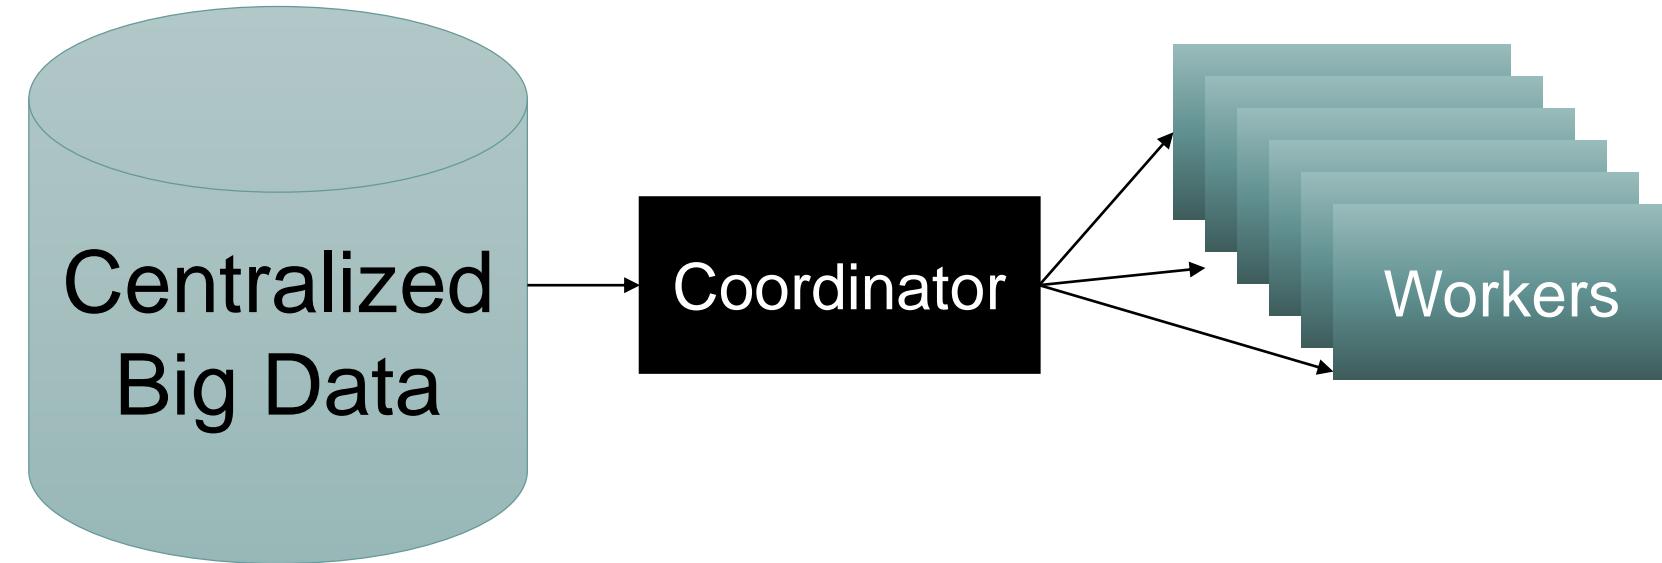
**Distributed Computing**

Big Data Storage

Cloud Services

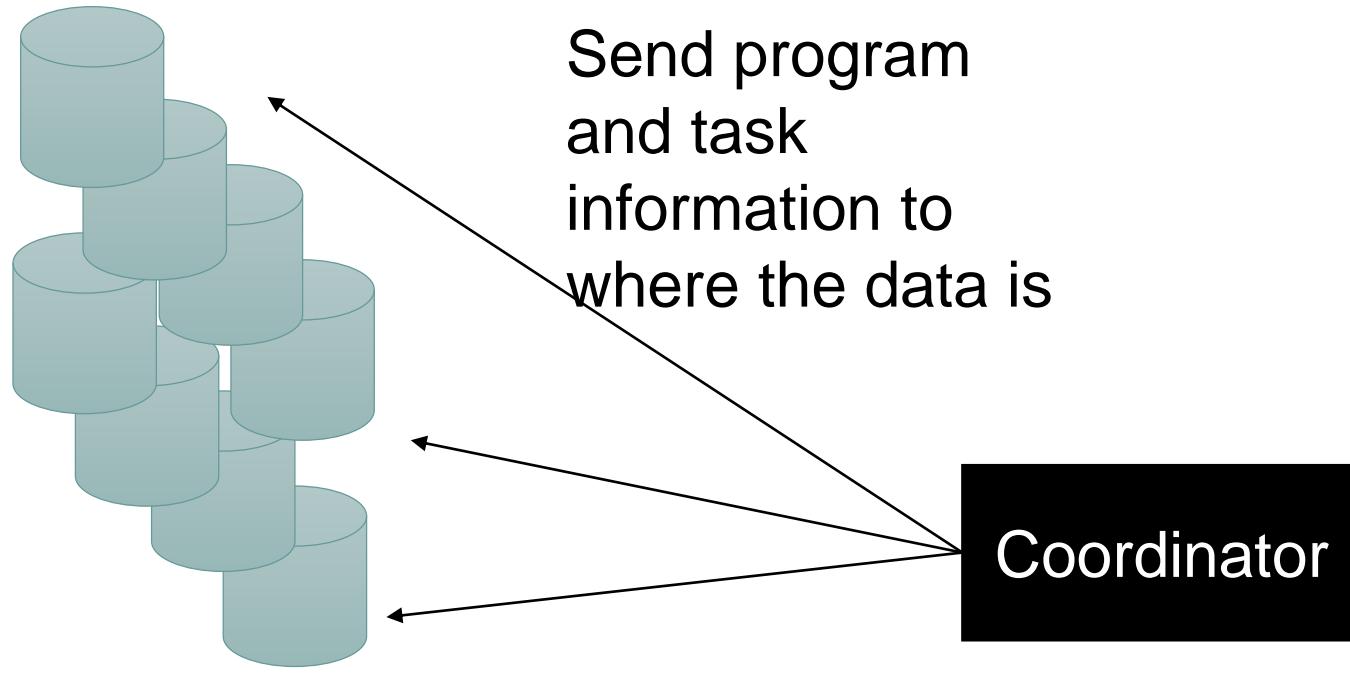
Coordination/  
Cluster  
Management

# Traditional Distributed Computing



Ship data to computation paradigm  
e.g., High performance computing (HPC)

# Big-data Computing



Storage/Compute  
Nodes

Ship compute to data paradigm

# MapReduce



- › Invented by Google
- › Implemented in the open-source Hadoop system
- › A functional programming technique
- › The program is expressed in two functions, Map and Reduce
- › Each function processes one or a few records at a time
- › The functions are applied in parallel to trillions of records

# Example: Word Count



if you cannot fly,  
then run.  
If you cannot  
run, then walk.  
If you cannot  
walk, then crawl  
But whatever  
you do you  
have to keep  
moving forward

Input text file

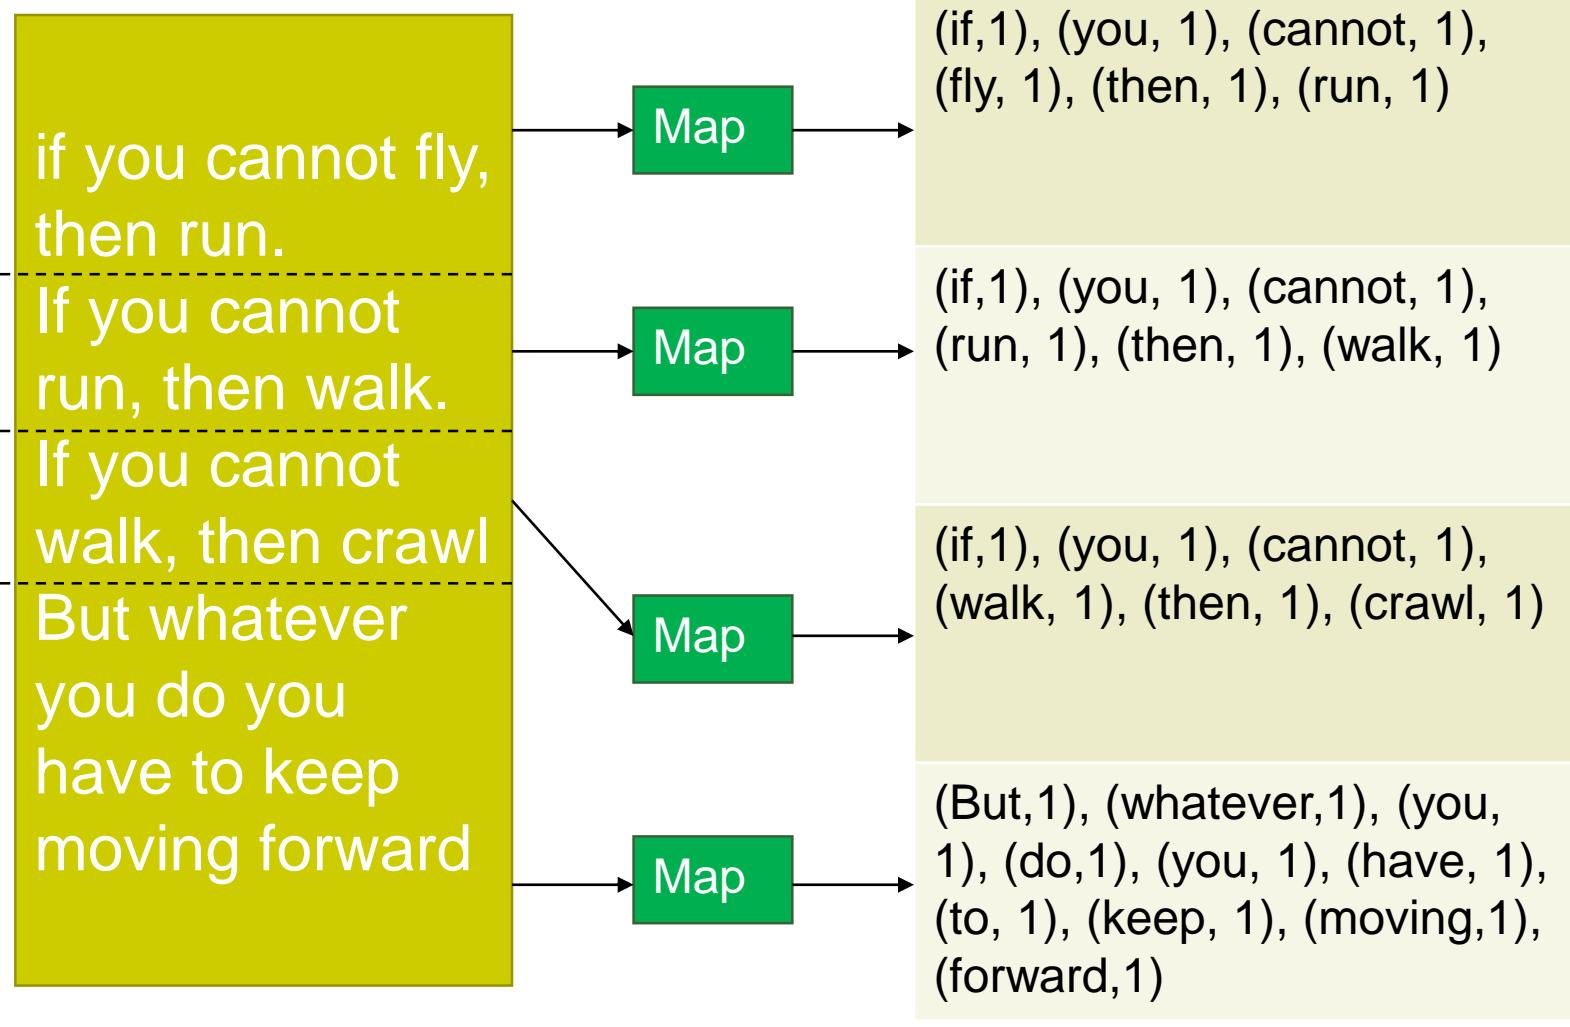
```
Map(line) {  
    split line into words  
    for each word w  
        output (w,1)  
}
```

```
Reduce(w, c[]) {  
    s = Sum(c)  
    output(w, s)  
}
```

you: 5  
cannot: 3  
walk: 2  
if: 3  
...

Output

# Example: Word Count (Map)



Input text file

Map output

# Example: Word Count (Shuffle)

(if,1), (you, 1), (cannot, 1),  
 (fly, 1), (then, 1), (run, 1)

(if,1), (you, 1), (cannot, 1),  
 (run, 1), (then, 1), (walk, 1)

(if,1), (you, 1), (cannot, 1),  
 (walk, 1), (then, 1), (crawl, 1)

(But,1), (whatever,1), (you,  
 1), (do,1), (you, 1), (have, 1),  
 (to, 1), (keep, 1), (moving,1),  
 (forward,1)

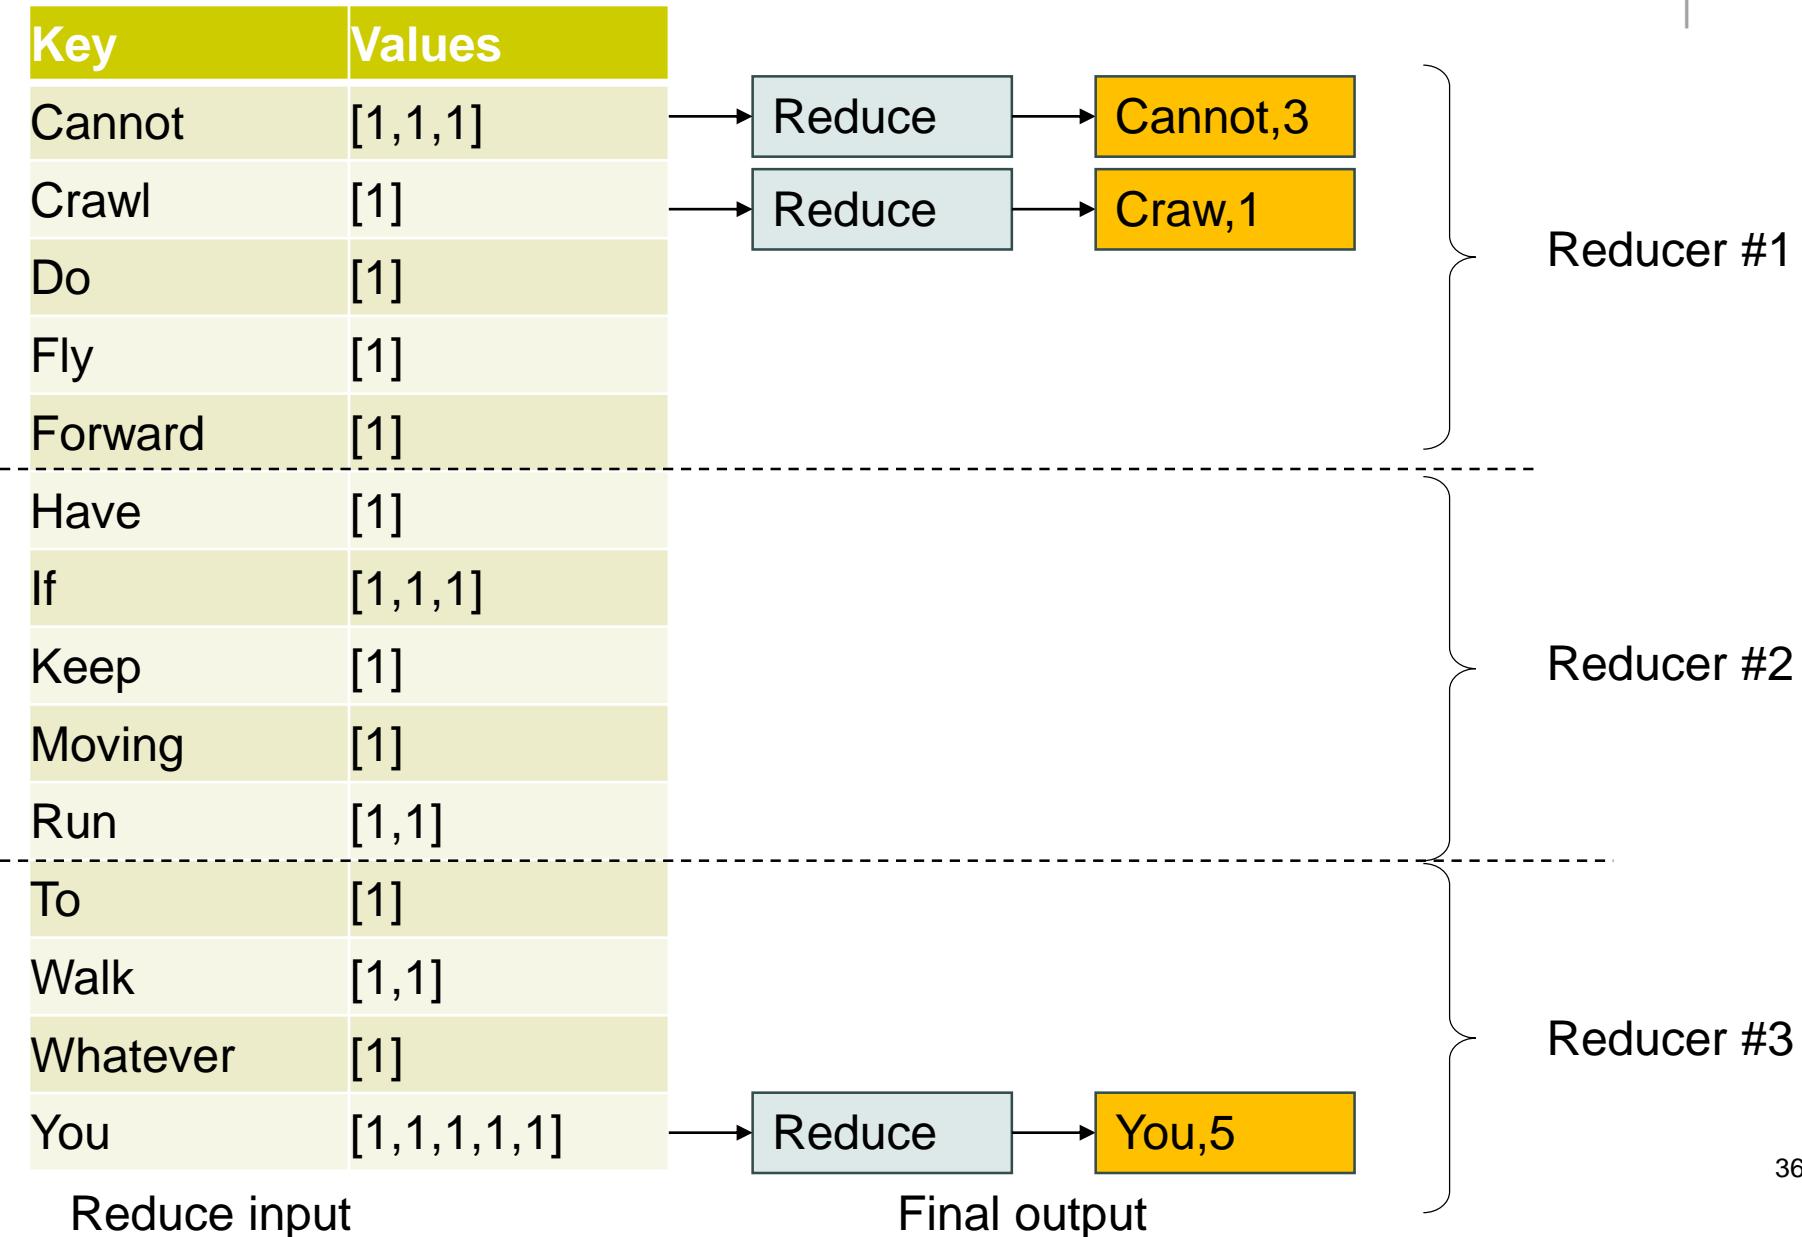
Map output

Group by  
key



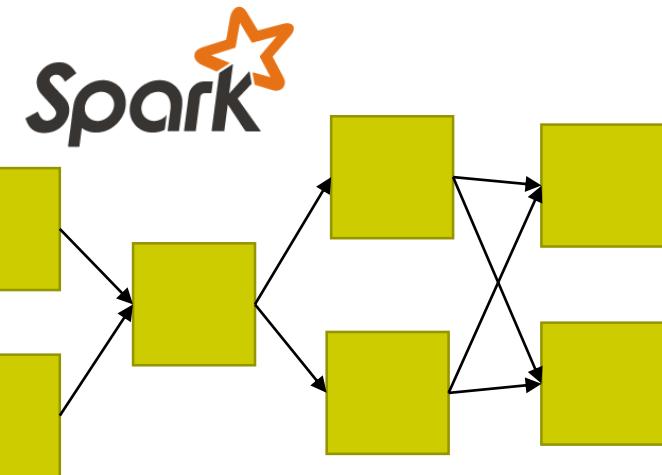
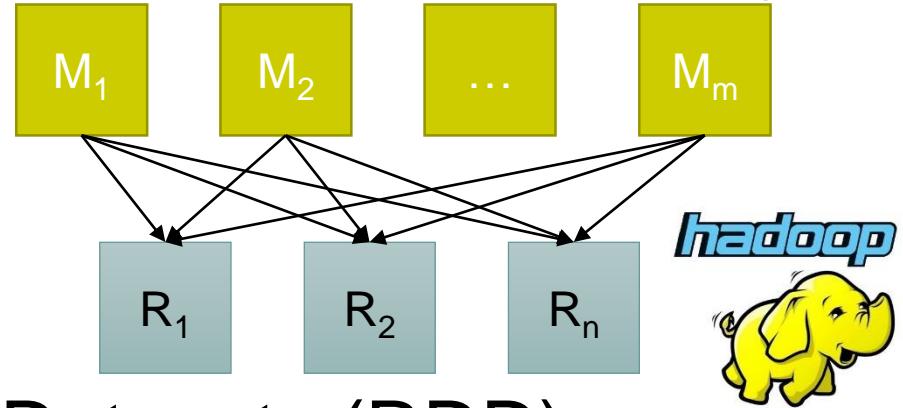
Key	Values
Cannot	[1,1,1]
Crawl	[1]
Do	[1]
Fly	[1]
Forward	[1]
Have	[1]
If	[1,1,1]
Keep	[1]
Moving	[1]
Run	[1,1]
To	[1]
Walk	[1,1]
Whatever	[1]
You	[1,1,1,1,1]

# Example: Word Count (Reduce)



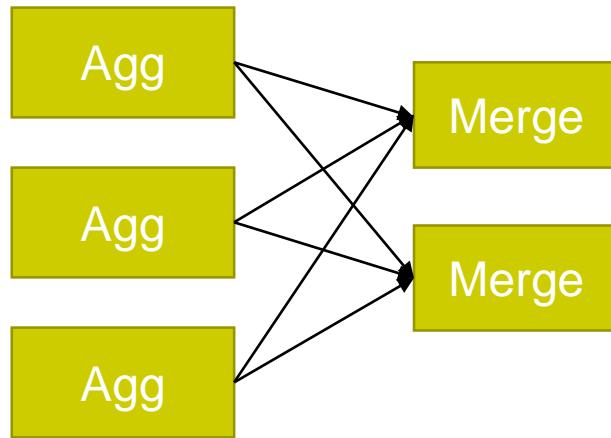
# Task Execution

- › MapReduce
  - › Map-Shuffle- Reduce
  - › Resiliency through materialization
- › Resilient Distributed Datasets (RDD)
- › Directed-Acyclic-Graph (DAG)
- › In-memory processing
- › Resiliency through lineages
- › Hyracks
- › AsterixDB
- › Stragglers
- › Load balance

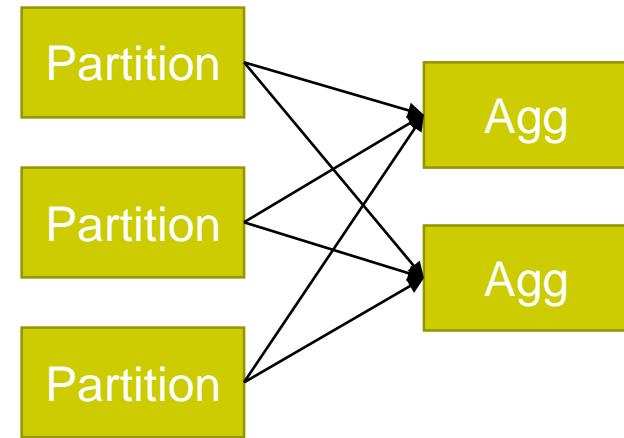


# Query Optimization

- › Finding the most efficient query plan
- › e.g., grouped aggregation



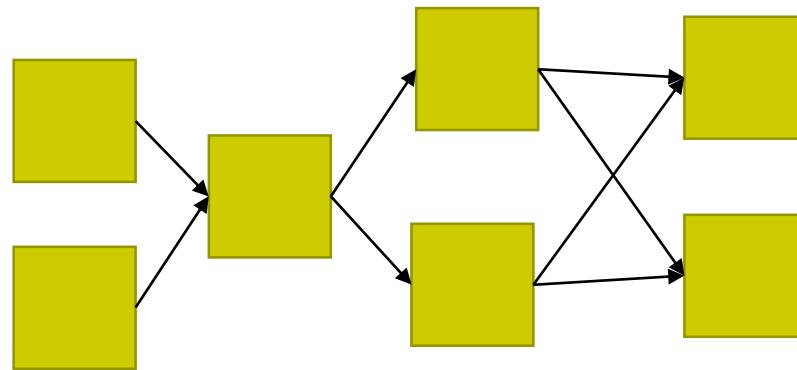
Vs



- › Cost model (CPU – Disk – Network)

# Provenance

- › Debugging in distributed systems is painful



- › We need to keep track of transformations on each record

# High-level Languages



Big-data Libraries

High-level Languages

Distributed Computing

Big Data Storage

Cloud Services

Coordination/  
Cluster  
Management

# Declarative MapReduce



- › MapReduce has been used to create many reusable operators (e.g., relational operators)
- › Filter Map
- › Aggregate Map Reduce
- › Grouped aggregated Map Reduce
- › Equi-join Map Reduce
- › Non-equijoin Map Reduce

# Declarative Languages



- › Describe what you want to do not how to do it
- › The most popular example is SQL
- › Can we compile SQL queries into MapReduce program(s)?

# Pig

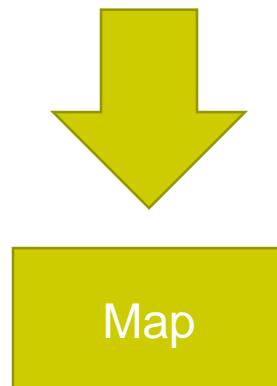


- › A system built on-top of Hadoop (Now supports Spark as well)
- › Provides a SQL-ETL-like query language termed Pig Latin
- › Compiles Pig Latin programs into MapReduce programs

# Examples

- › Filter: Return all the lines that have a user-specified response code, e.g., 200.

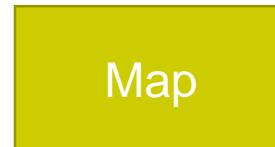
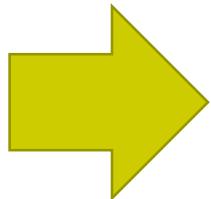
```
log = LOAD 'logs.csv' USING PigStorage()  
    AS (host, time, method, url, response, bytes);  
ok_lines = FILTER log BY response = '200';  
STORE ok_lines into 'filtered_output';
```



# Examples

- › Grouped aggregate
- › Find the total number of bytes per response code

```
log = LOAD 'logs.csv' USING PigStorage()
    AS (host, time, method, url, response, bytes: int);
grouped = GROUP log BY response;
grouped_aggregate =
    FOREACH grouped GENERATE group, SUM(bytes);
STORE grouped_aggregate into 'grouped_output';
```



# Examples

- › Grouped aggregate
- › Find the **average** number of bytes per response code

```
log = LOAD 'logs.csv' USING PigStorage()
    AS (host, time, method, url, response, bytes: int);
grouped = GROUP log BY response;
grouped_aggregate =
    FOREACH grouped GENERATE group, AVG(bytes);
STORE grouped_aggregate into 'grouped_output';
```

# Examples

- Join: Find pairs of requests that ask for the **same URL**, coming from the **same source**

```
log1 = LOAD 'logs.csv' USING PigStorage()  
    AS (host, time, method, url, response, bytes: int);  
log2 = LOAD 'logs.csv' USING PigStorage()  
    AS (host, time, method, url, response, bytes: int);  
joined = JOIN log1 BY (url, host),  
        log2 BY (url, host);
```

# Examples

- Join: Find pairs of requests that ask for the **same URL**, coming from the **same source** and happened **within an hour of each other**

```
log1 = LOAD 'logs.csv' USING PigStorage()  
    AS (host, time, method, url, response, bytes: int);  
log2 = LOAD 'logs.csv' USING PigStorage()  
    AS (host, time, method, url, response, bytes: int);  
joined = JOIN log1 BY (url, host),  
        log2 BY (url, host);  
filtered = FILTER joined BY  
    ABS(log1::time - log2::time) < 3600000;
```

# Additional Features



- › Lazy execution
  - › Nothing gets actually executed until the STORE command is reached
- › Consolidation of map-only jobs
  - › Map-only jobs (FILTER and FOREACH) can be consolidated into a next job's map function or a previous job's reduce function

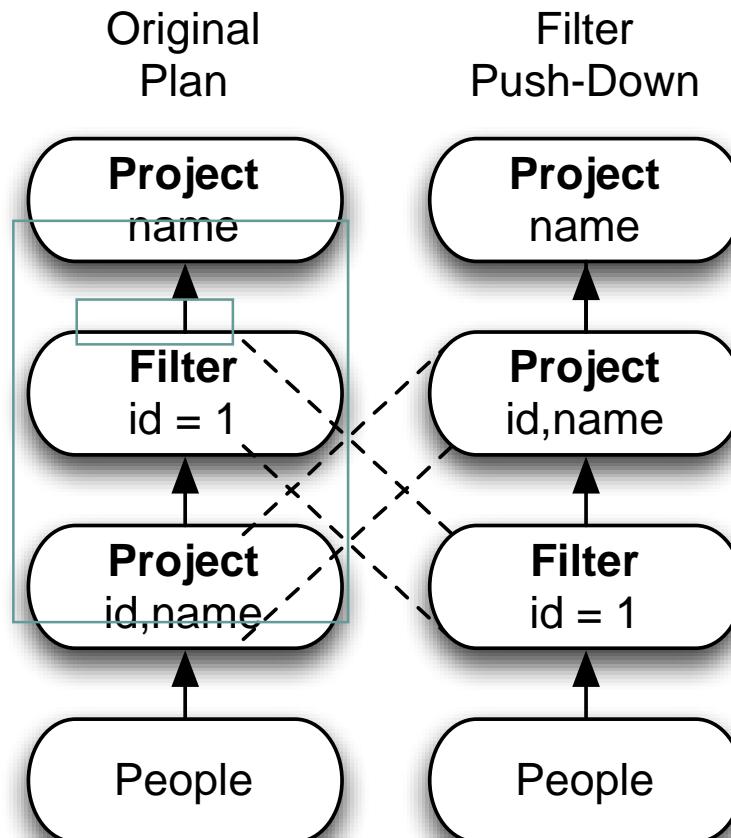
# SparkSQL

- › Redesigned to consider Spark query model
- › Supports all the popular relational operators
- › Can be intermixed with RDD operations
- › Uses the Dataframe API as an enhancement to the RDD API

Dataframe = RDD + schema

# Catalyst Query Optimizer

- › Extensible rule-based optimizer
- › Users can define their own rules



# Big-data Libraries



**Big-data Libraries**

High-level Languages

Distributed Computing

Big Data Storage

Cloud Services

Coordination/  
Cluster  
Management

# Specialized Big Data Systems



- › The flexible big-data systems allowed developers to build many high-level applications
- › Graph processing
  - › Giraph
  - › GraphX
- › Machine Learning
  - › Mahout
  - › MLlib

# Cloud Services



Big-data Libraries

High-level Languages

Distributed Computing

Big Data Storage

Cloud Services

Coordination/  
Cluster  
Management

# Elastic Clouds



- › Clusters are expensive to buy and manage
- › Machine could be idle most of the time
- › Elastic clouds enables creation and destruction of on-demand clusters



Google  
Cloud Platform



Microsoft  
Azure

# Big-data on the Cloud



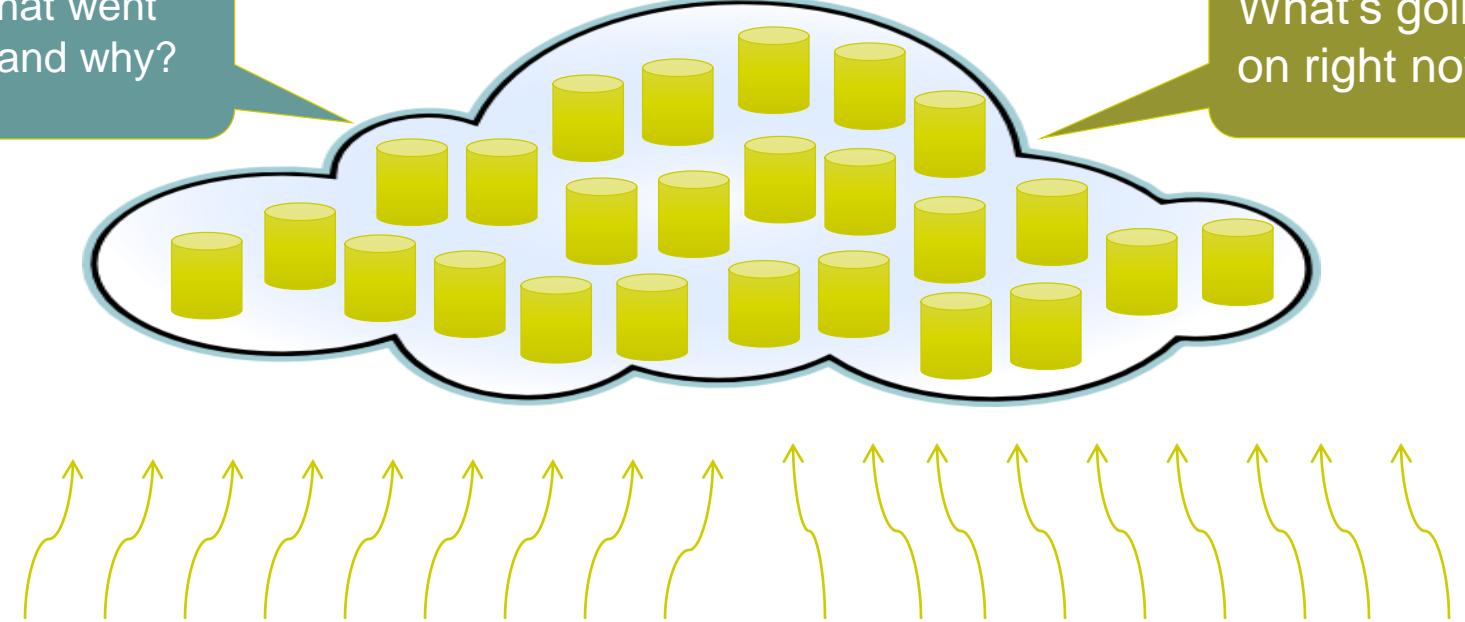
- Most cloud services provide ready-to-go big-data clusters, e.g., Hadoop or Spark
- You can easily tailor the hardware specifications according to your application needs, e.g., storage, memory, computing, and network bandwidth
- You only pay for the hours where the machines are running
- Make sure to terminate the machines after you are done!!!

# Big Data / Web Warehousing

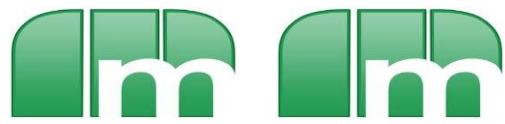


So what went  
on – and why?

What's going  
on right now?



# Also: Today's Big Data Tangle

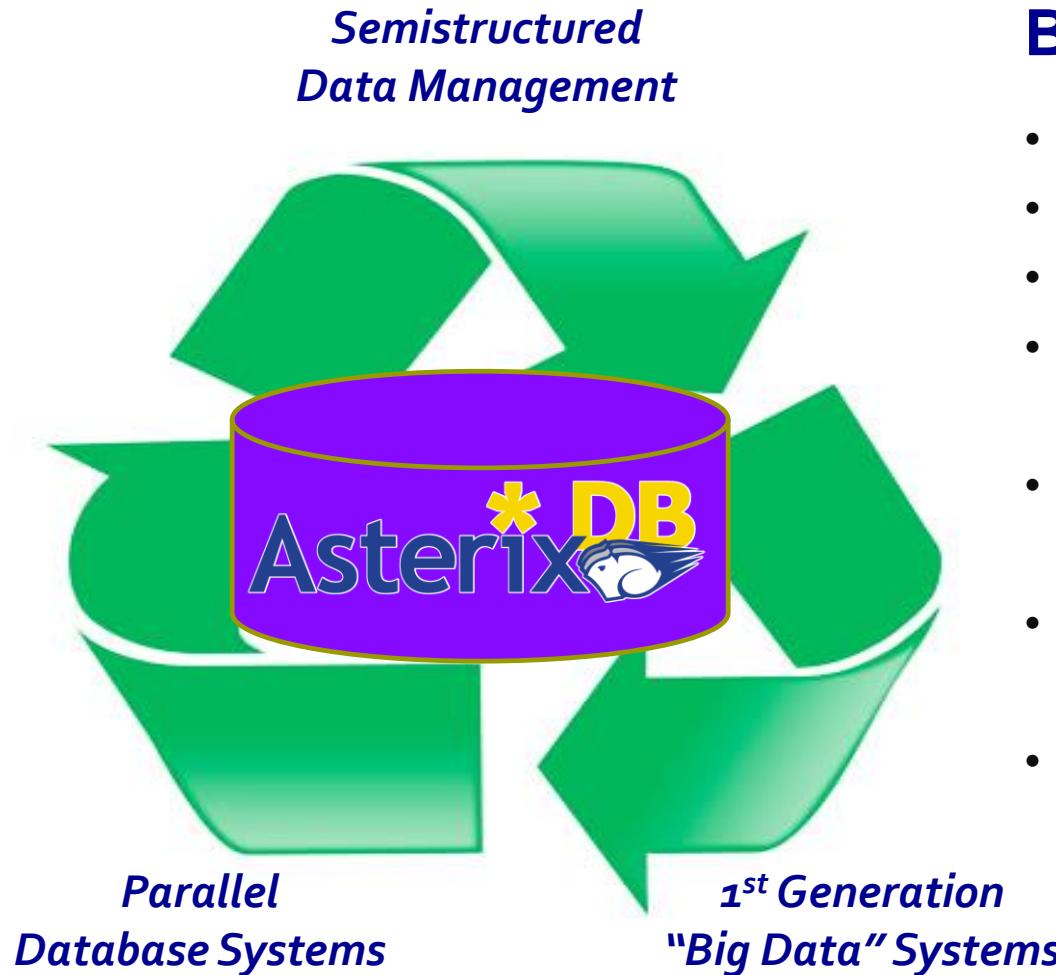


*Cassandra*

Scalable, eventually consistent, distributed, structured key-value store.



# AsterixDB: “One Size Fits a Bunch”



## BDMS Desiderata:

- Able to **manage** data
- **Flexible** data model
- Full **query** capability
- Continuous data **ingestion**
- Efficient and robust **parallel** runtime
- Cost **proportional** to task at hand
- Support “**Big Data**” data types”
- 
- 
-

# ASTERIX Data Model (ADM)

```
CREATE DATAVERSE TinySocial;
```

```
USE TinySocial;
```

```
CREATE TYPE GleambookUserType AS {
```

```
    id: int,  
    alias: string,  
    name: string,  
    userSince: datetime,  
    friendIds: {{ int }},  
    employment: [EmploymentType]  
};
```

```
CREATE TYPE EmploymentType AS {
```

```
    organizationName: string,  
    startDate: date,  
    endDate: date?
```

```
CREATE DATASET GleambookUsers  
(GleambookUserType)
```

```
PRIMARY KEY id;
```

## *Highlights include:*

- › JSON++ based data model
- › Rich type support (spatial, temporal, ...)
- › Records, lists, bags
- › *Open vs. closed types*

# Other DDL Features

```

CREATE INDEX gbUserSincelidx ON GleambookUsers(userSince);
CREATE INDEX gbAuthorIdx ON GleambookMessages(authorId) TYPE BTREE;
CREATE INDEX gbSenderLocIndex ON GleambookMessages(senderLocation) TYPE RTREE;
CREATE INDEX gbMessageIdx ON GleambookMessages(message) TYPE KEYWORD;
//----- and also -----

```

**CREATE TYPE** AccessLogType **AS CLOSED**

```

{ ip: string, time: string, user: string, verb: string, `path`: string, stat: int32, size: int32 };
CREATE EXTERNAL DATASET AccessLog(AccessLogType) USING localfs
(("path"="localhost:///Users/mikejcarey/extdemo/accesses.txt"),
("format"="delimited-text"), ("delimiter"="|"));

```

**CREATE FEED** myMsgFeed **USING** socket\_adapter

```

(("sockets"="127.0.0.1:10001"), ("address-type"="IP"),
("type-name"="GleambookMessageType"), ("format"="adm"));

```

**CONNECT FEED** myMsgFeed **TO DATASET** GleambookMessages;

**START FEED** myMsgFeed;

*External data highlights:*

- Equal opportunity access
- Feeds to “keep everything!”
- Ingestion, *not streams*<sup>61</sup>

# ASTERIX Queries (SQL++ or AQL)

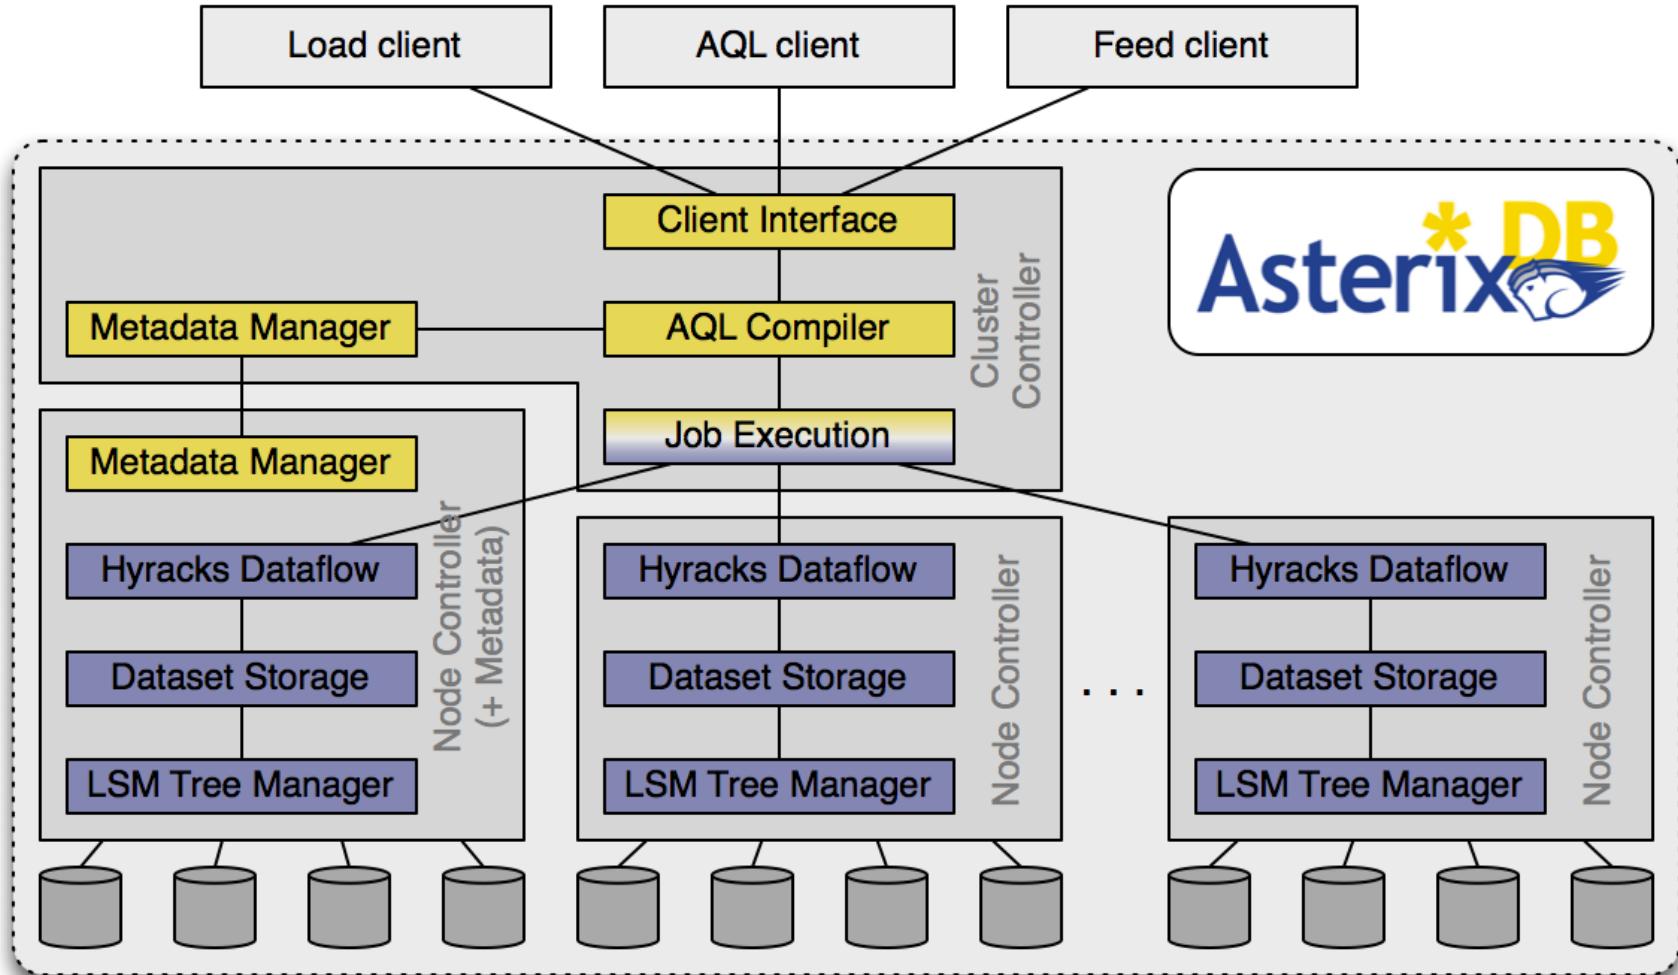


- Q1: List the user names and messages sent by Gleambook social network users with less than 3 friends:

```
SELECT user.name AS uname,  
       (SELECT VALUE msg.message  
        FROM GleambookMessages msg  
        WHERE msg.authorId = user.id) AS messages  
FROM GleambookUsers user  
WHERE COLL_COUNT(user.friendIds) < 3;
```

```
{ "uname": "NilaMilliron", "messages": [ ] }  
{ "uname": "WoodrowNehling", "messages": [ " love acast its 3G is good:") ] }  
{ "uname": "IsbelDull", "messages": [ " like product-y the plan is amazing", " like  
product-z its platform is mind-blowing" ] }  
...
```

# AsterixDB System Overview





# Hyracks Dataflow Runtime

- Partitioned-parallel platform for data-intensive computing
- Job = dataflow DAG of operators and connectors
  - Operators consume and produce **partitions** of data
  - Connectors **route** (repartition) data between operators
- Hyracks vs. the “competition”
  - Based on time-tested parallel database principles
  - vs. Hadoop MR: More flexible model and less “pessimistic”
  - vs. newer SQL-on-Hadoop runtimes: Emphasis on out-of-core execution and adherence to memory budgets
  - Fast job activation, data pipelining, binary format, state-of-the-art DB style operators (hash-based, indexed, ...)
- Early test at Yahoo! Labs on 180 nodes (1440 cores, 720 disks)

# Hyracks (cont.)

## Query

```
use dataverse TinySocial

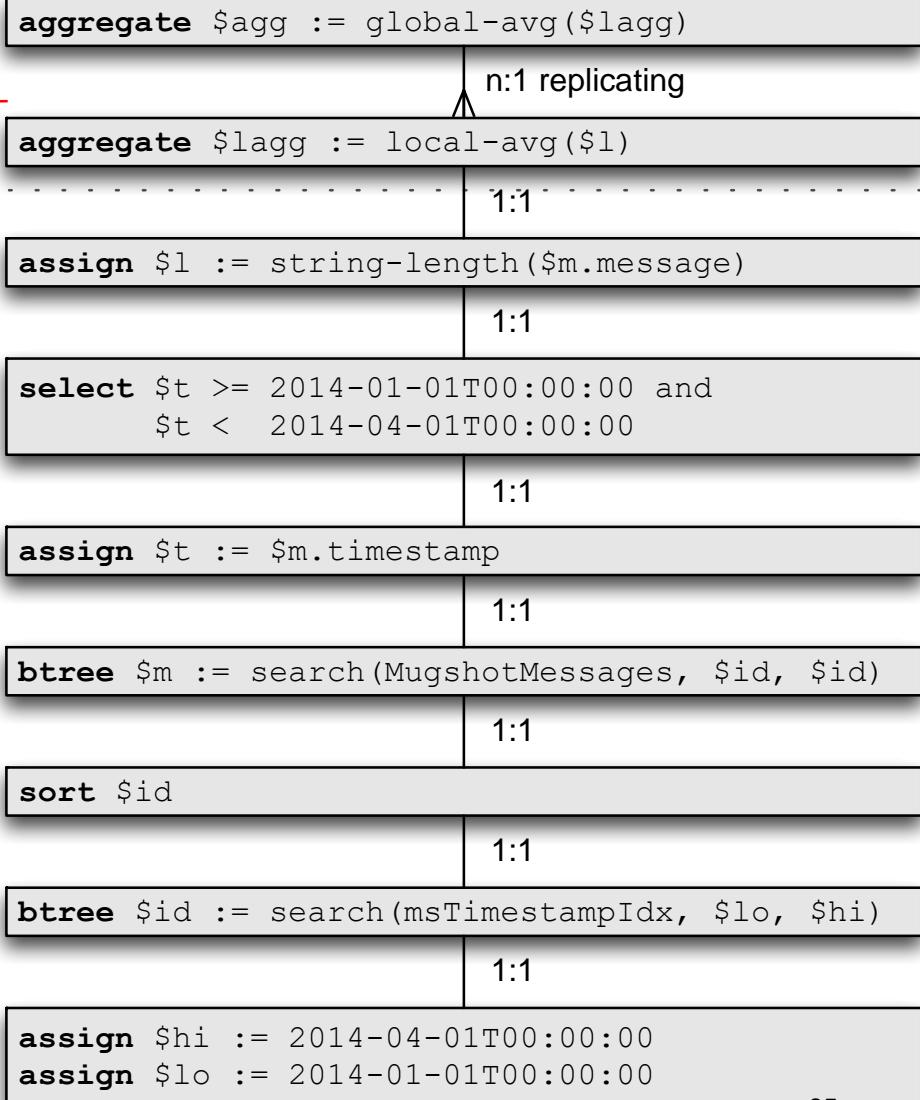
avg (
  for $m in dataset MugshotMessages
  where $m.timestamp >=
    datetime("2014-01-01T00:00:00")
    and $m.timestamp <
    datetime("2014-04-01T00:00:00")
  return string-length($m.message)
)}
```

Select Options   Clear Query   Run

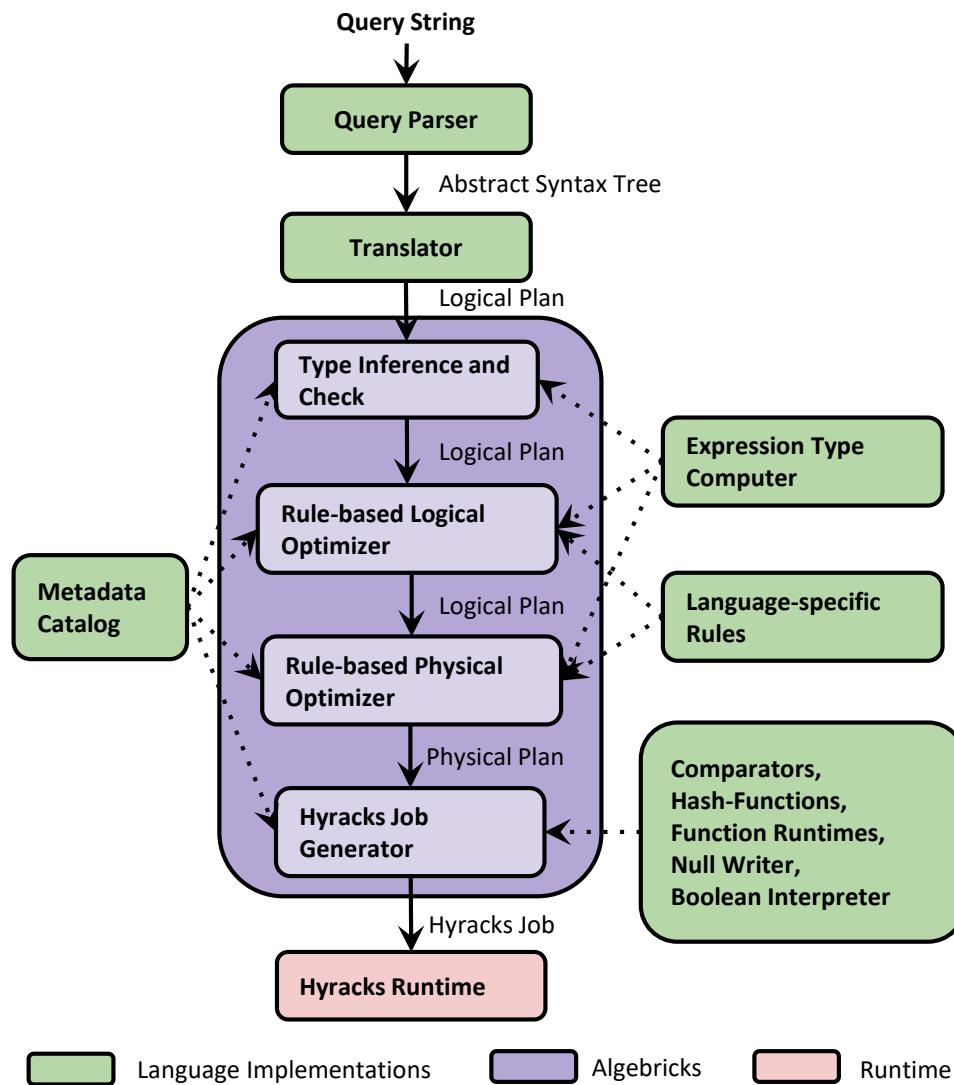
Print parsed expressions  
 Print rewritten expressions  
 Print logical plan  
 Print optimized logical plan  
 Print Hyracks job  
 Execute query

*Partitioned Parallelism!*

*Algebricks*



# Algebricks Query Compiler Framework



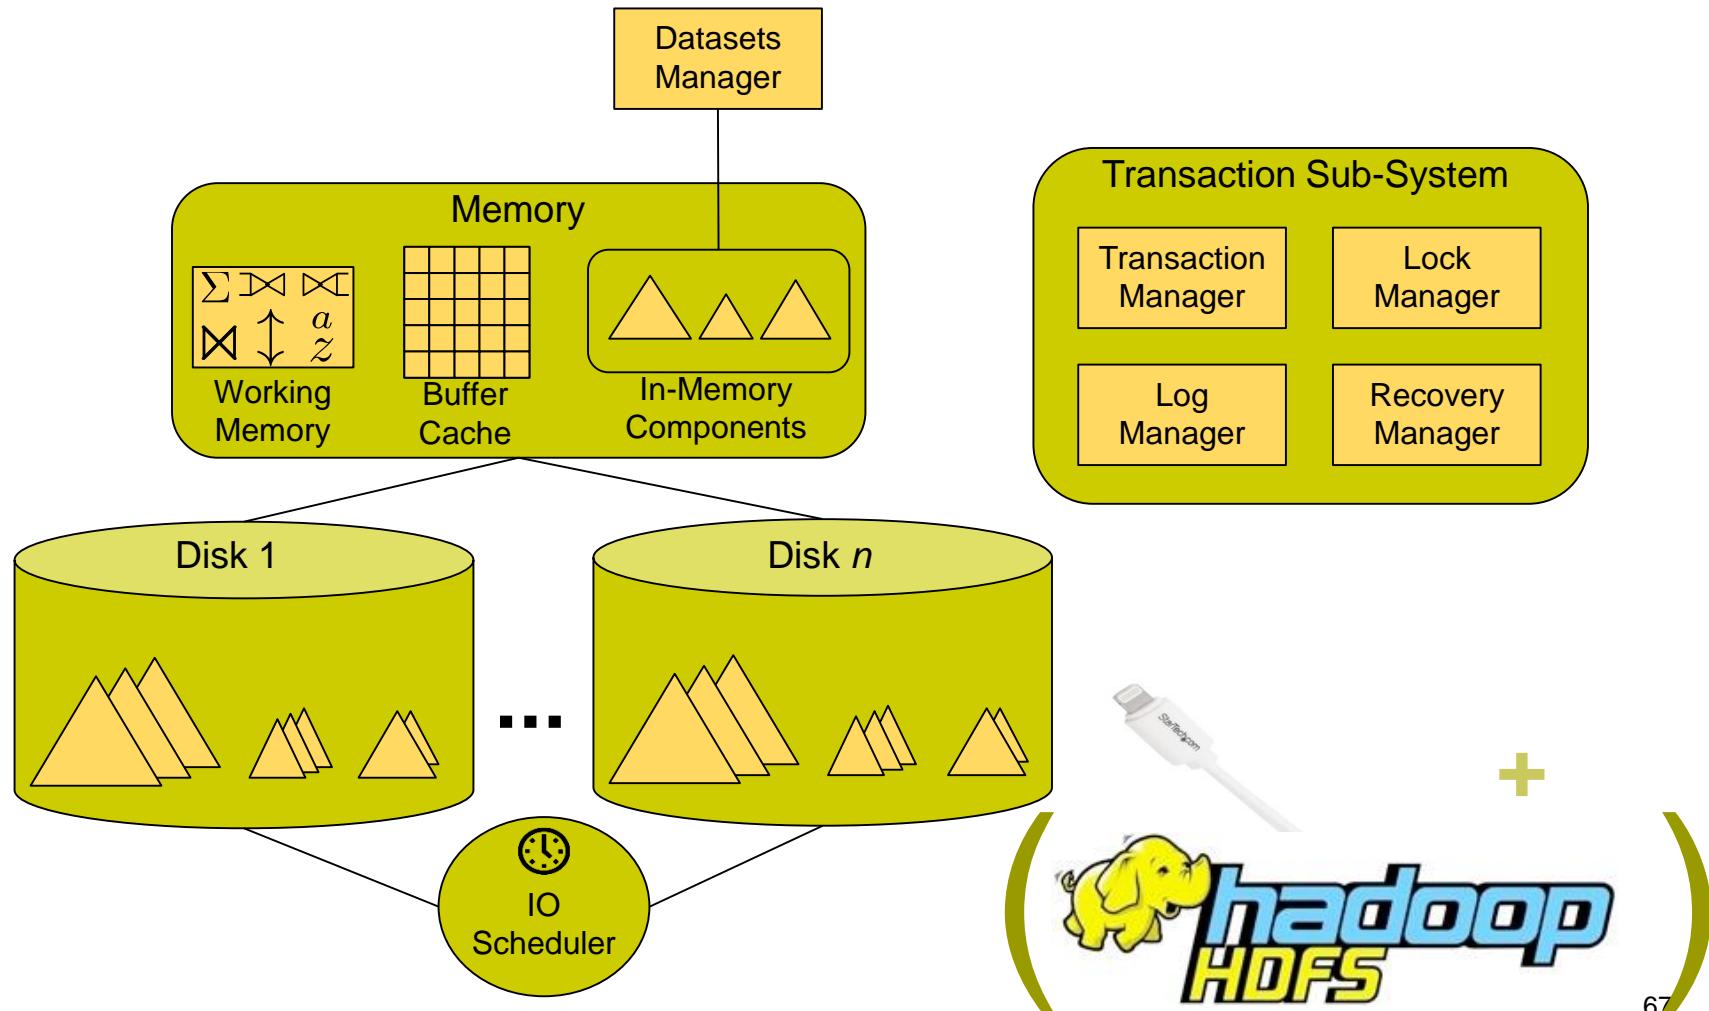
## Algebricks

- Logical Operators
- Logical Expressions
- Metadata Interface
- Model-Neutral Logical Rewrite Rules
- Physical Operators
- Model-Neutral Physical Rewrite Rules
- Hyracks Job Generator

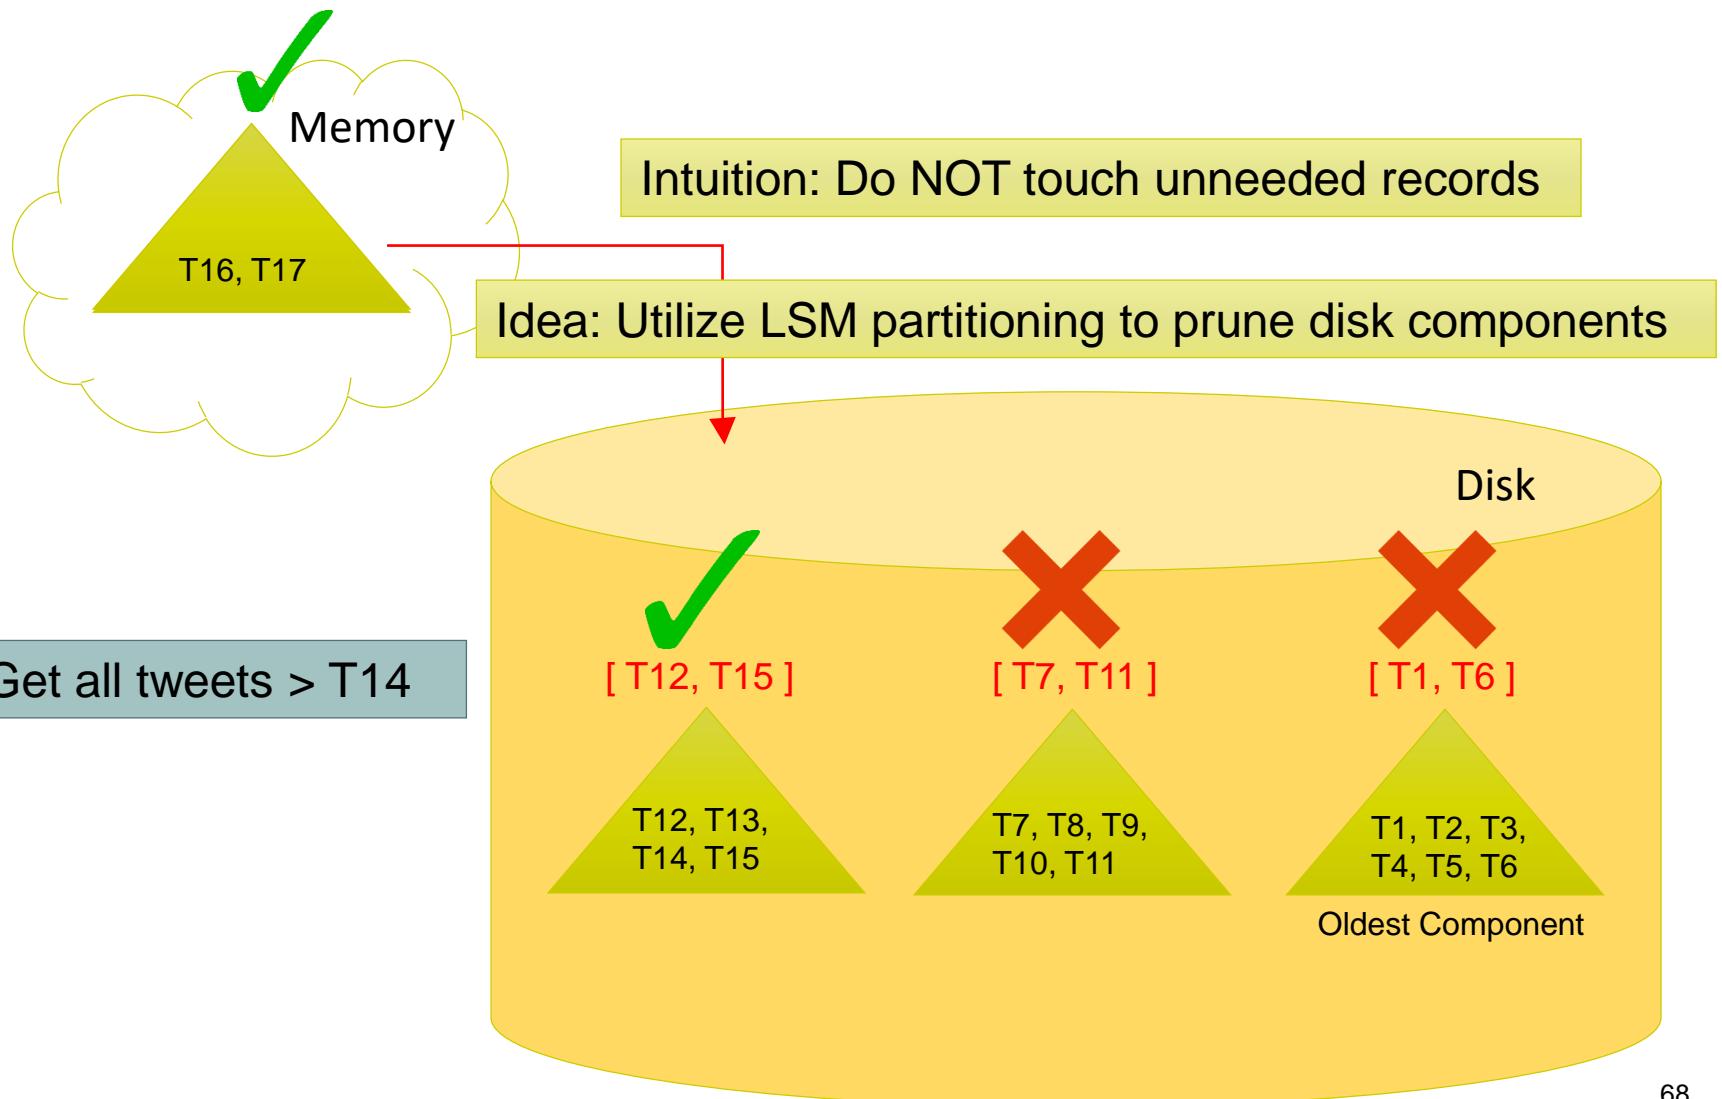
## Target Query Language

- Query Parser (AST)
- AST Translator
- Metadata Catalog
- Expression Type Computer
- Logical Rewrite Rules
- Physical Rewrite Rules
- Language Specifics

# Native Storage Management



# LSM-Based Filters



# For More Information



- › Apache AsterixDB home
  - › <http://asterixdb.apache.org/>
- › SQL++ Primer
  - › <http://asterixdb.apache.org/docs/0.9.4.1/sqlpp/primer-sqlpp.html>

# Conclusion



## Big-data Libraries

MLlib (Machine Learning), GraphX

## High-level Languages

SparkSQL, Pig, SQL++, HiveQL

## Distributed Computing

MapReduce (Hadoop and Google), Resilient Distributed Dataset (Spark), Hyracks (AsterixDB)

## Big Data Distributed Storage

Hadoop Distributed File System, Cloud storage systems (Amazon S3 and Google File System), Key-value stores

## Cloud Services

Amazon Web Services, Microsoft Azure, and Google Cloud Platform

## Coordination/Cluster Management

Oozie, Yarn, Kubernetes

# Thank You!

Questions?